

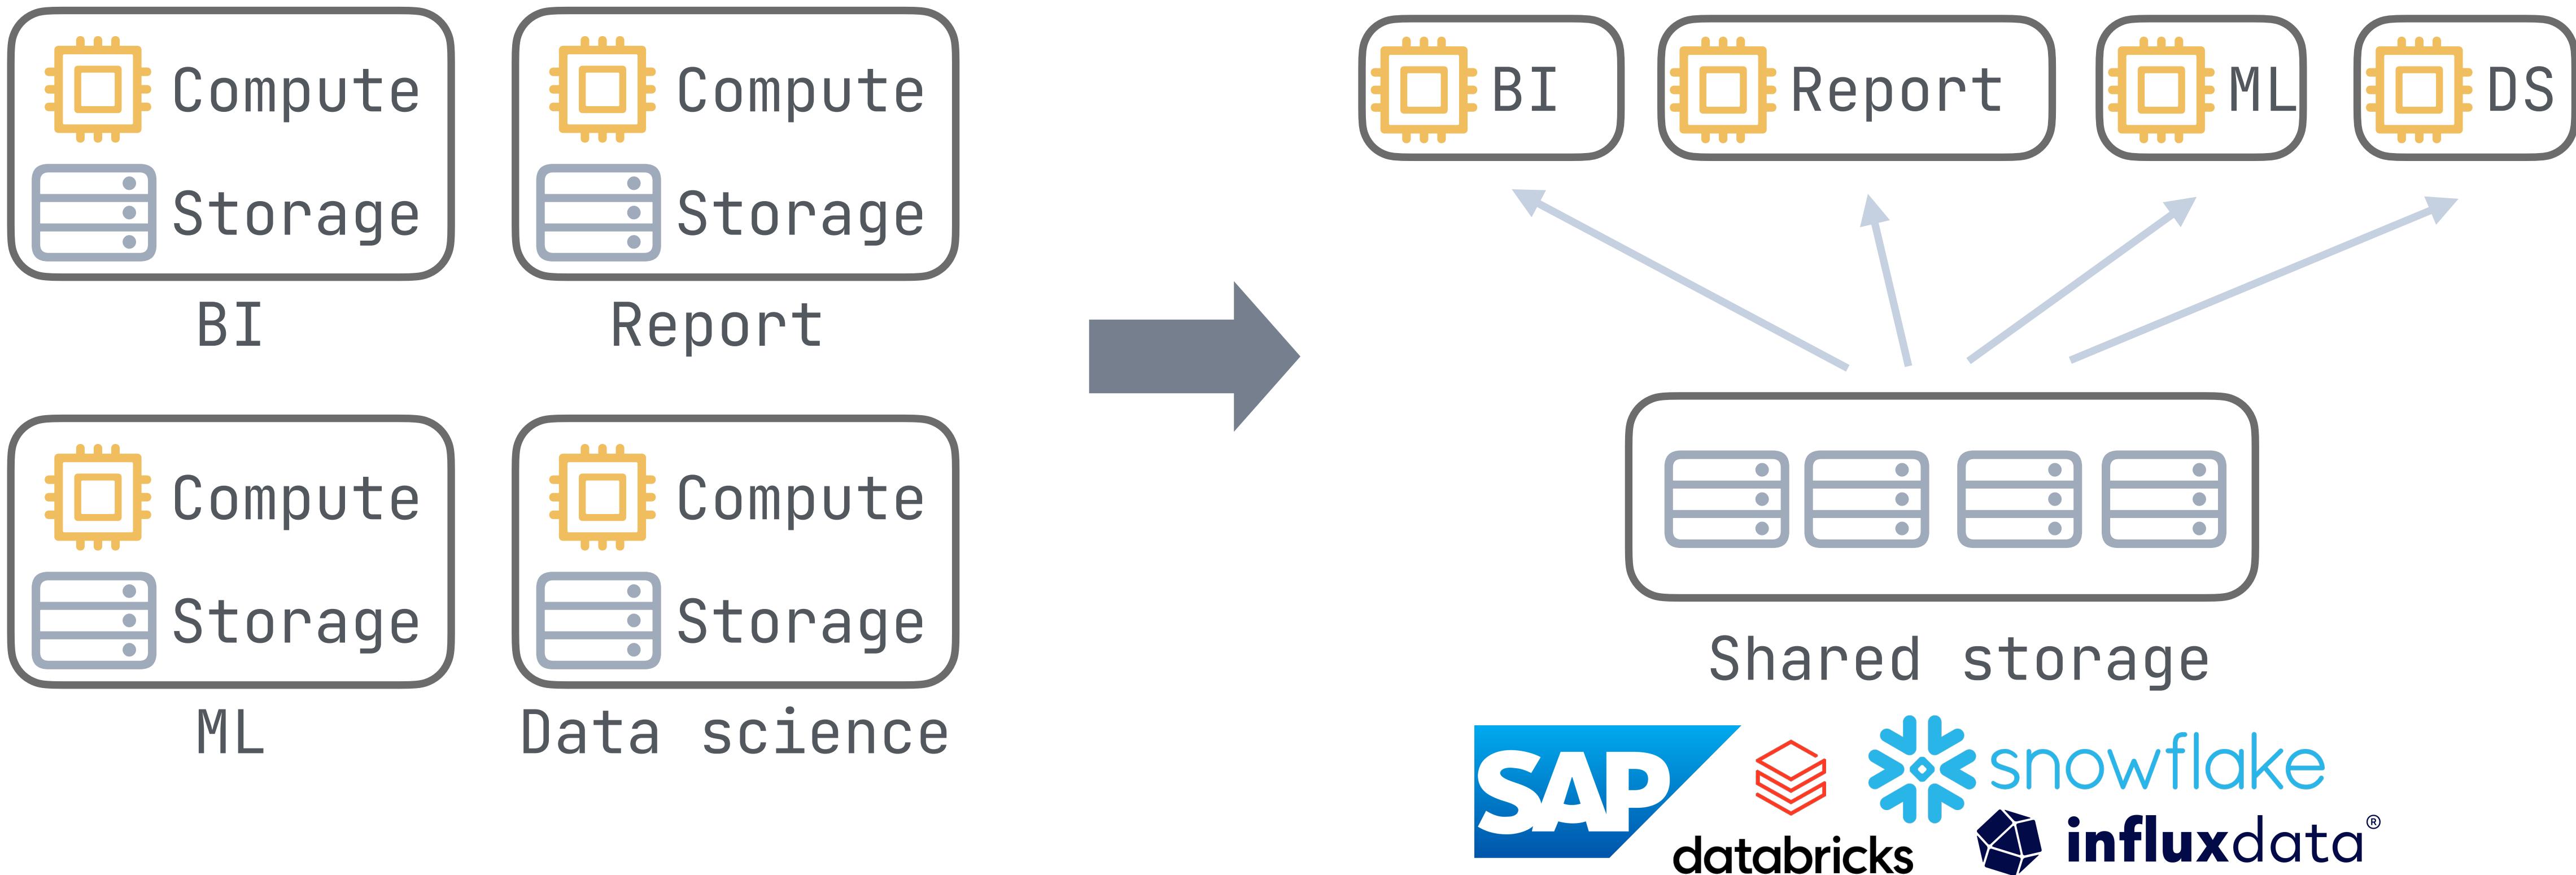
LiquidCache

Efficient Pushdown Caching for Cloud-Native Data Analytics

Xiangpeng Hao, University of Wisconsin-Madison

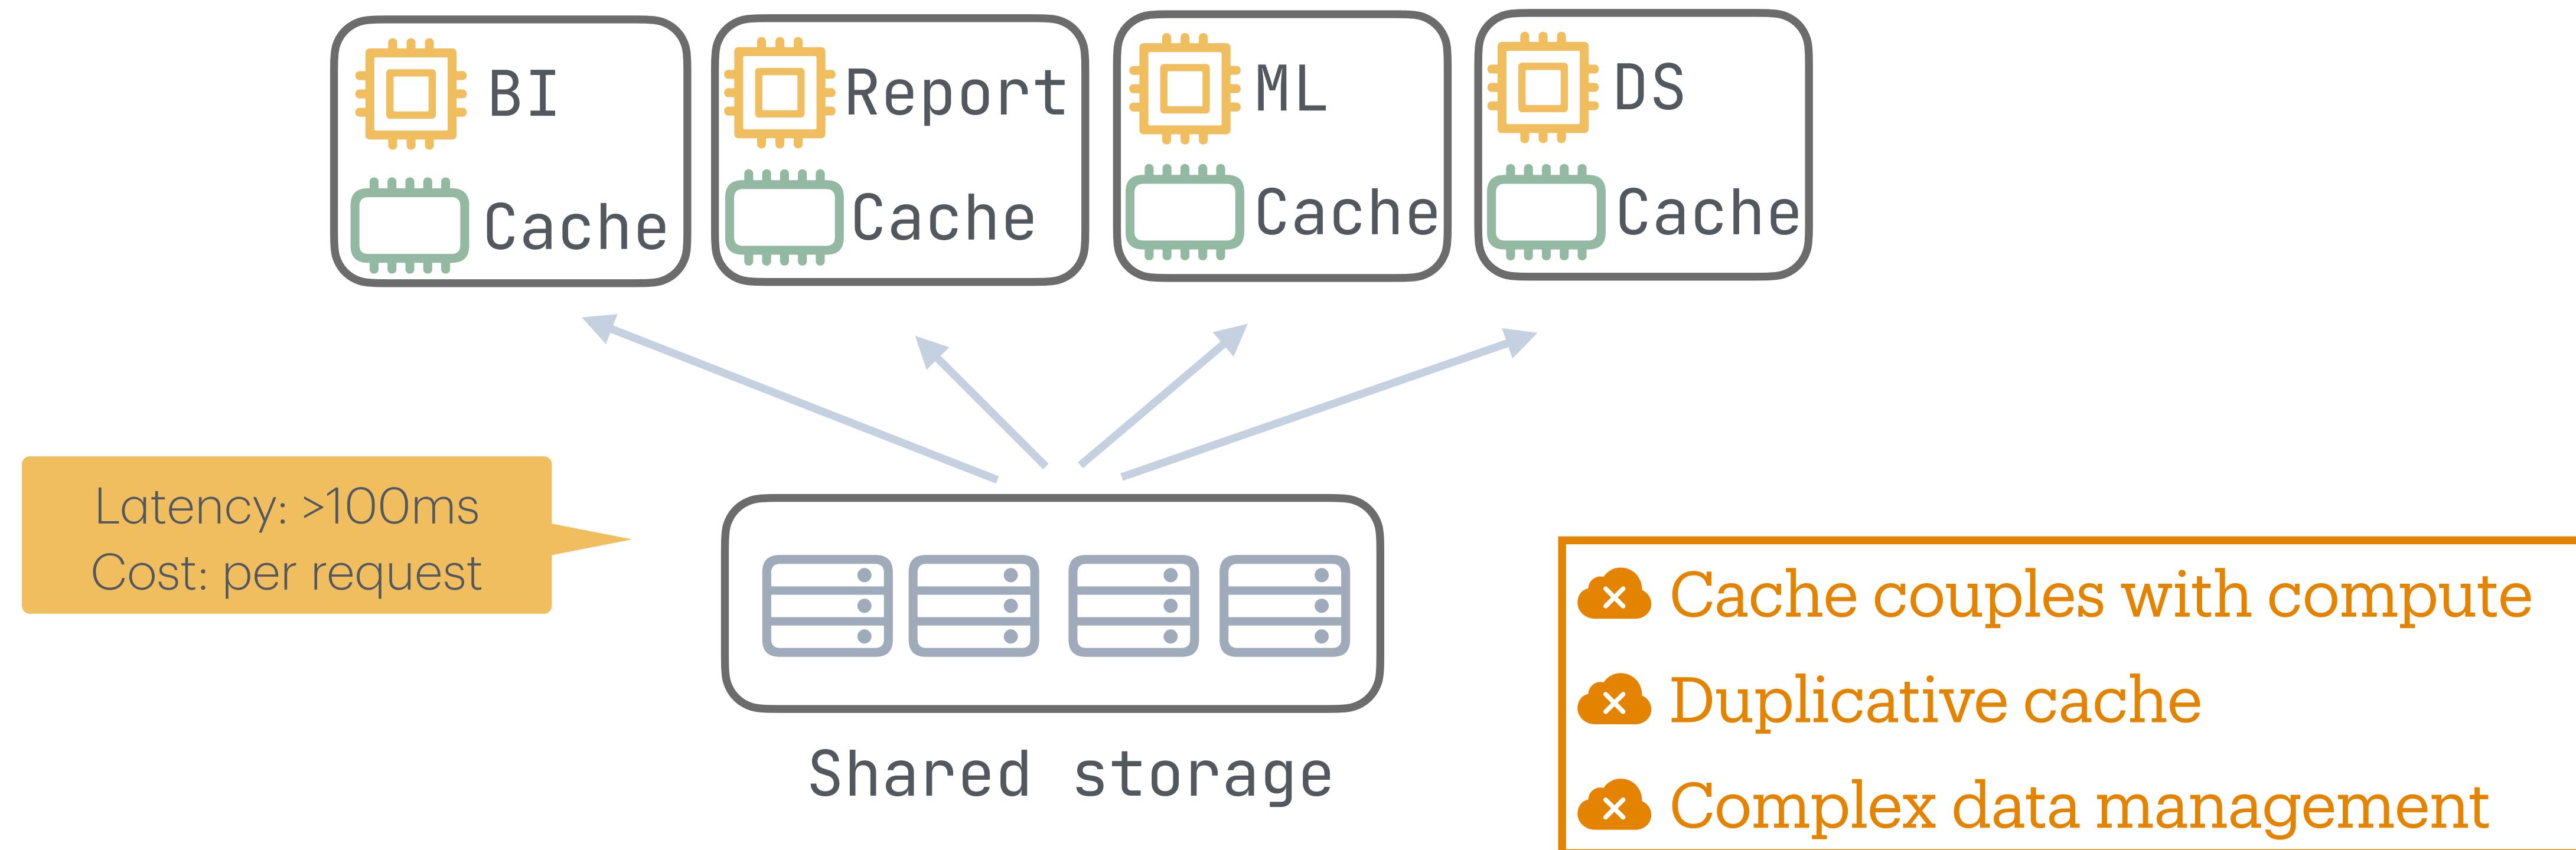
<https://xiangpeng.systems>

On premise → cloud (2010-2020)

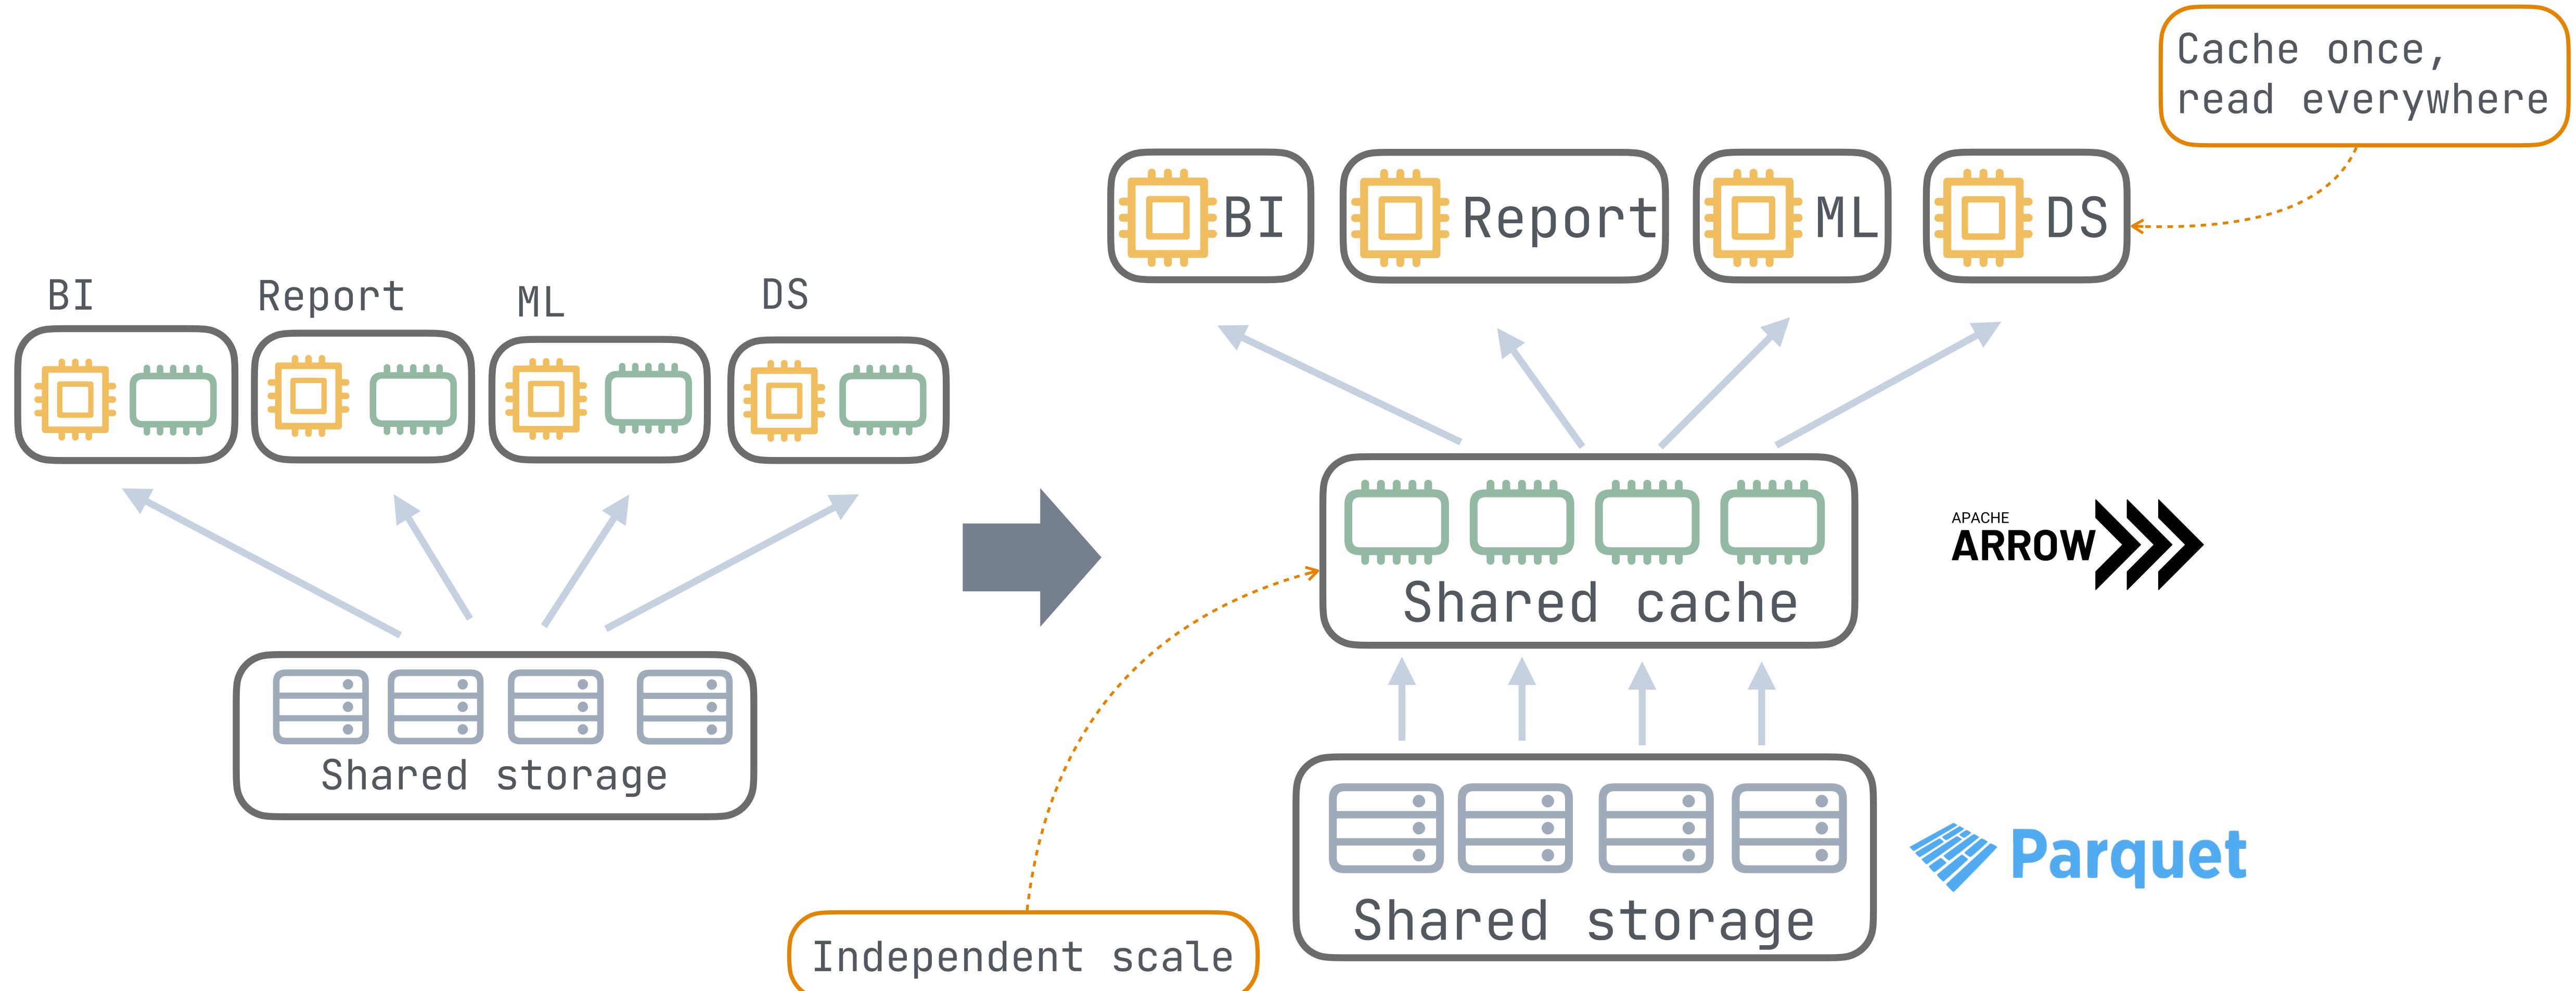


Elastic, durable, scalable

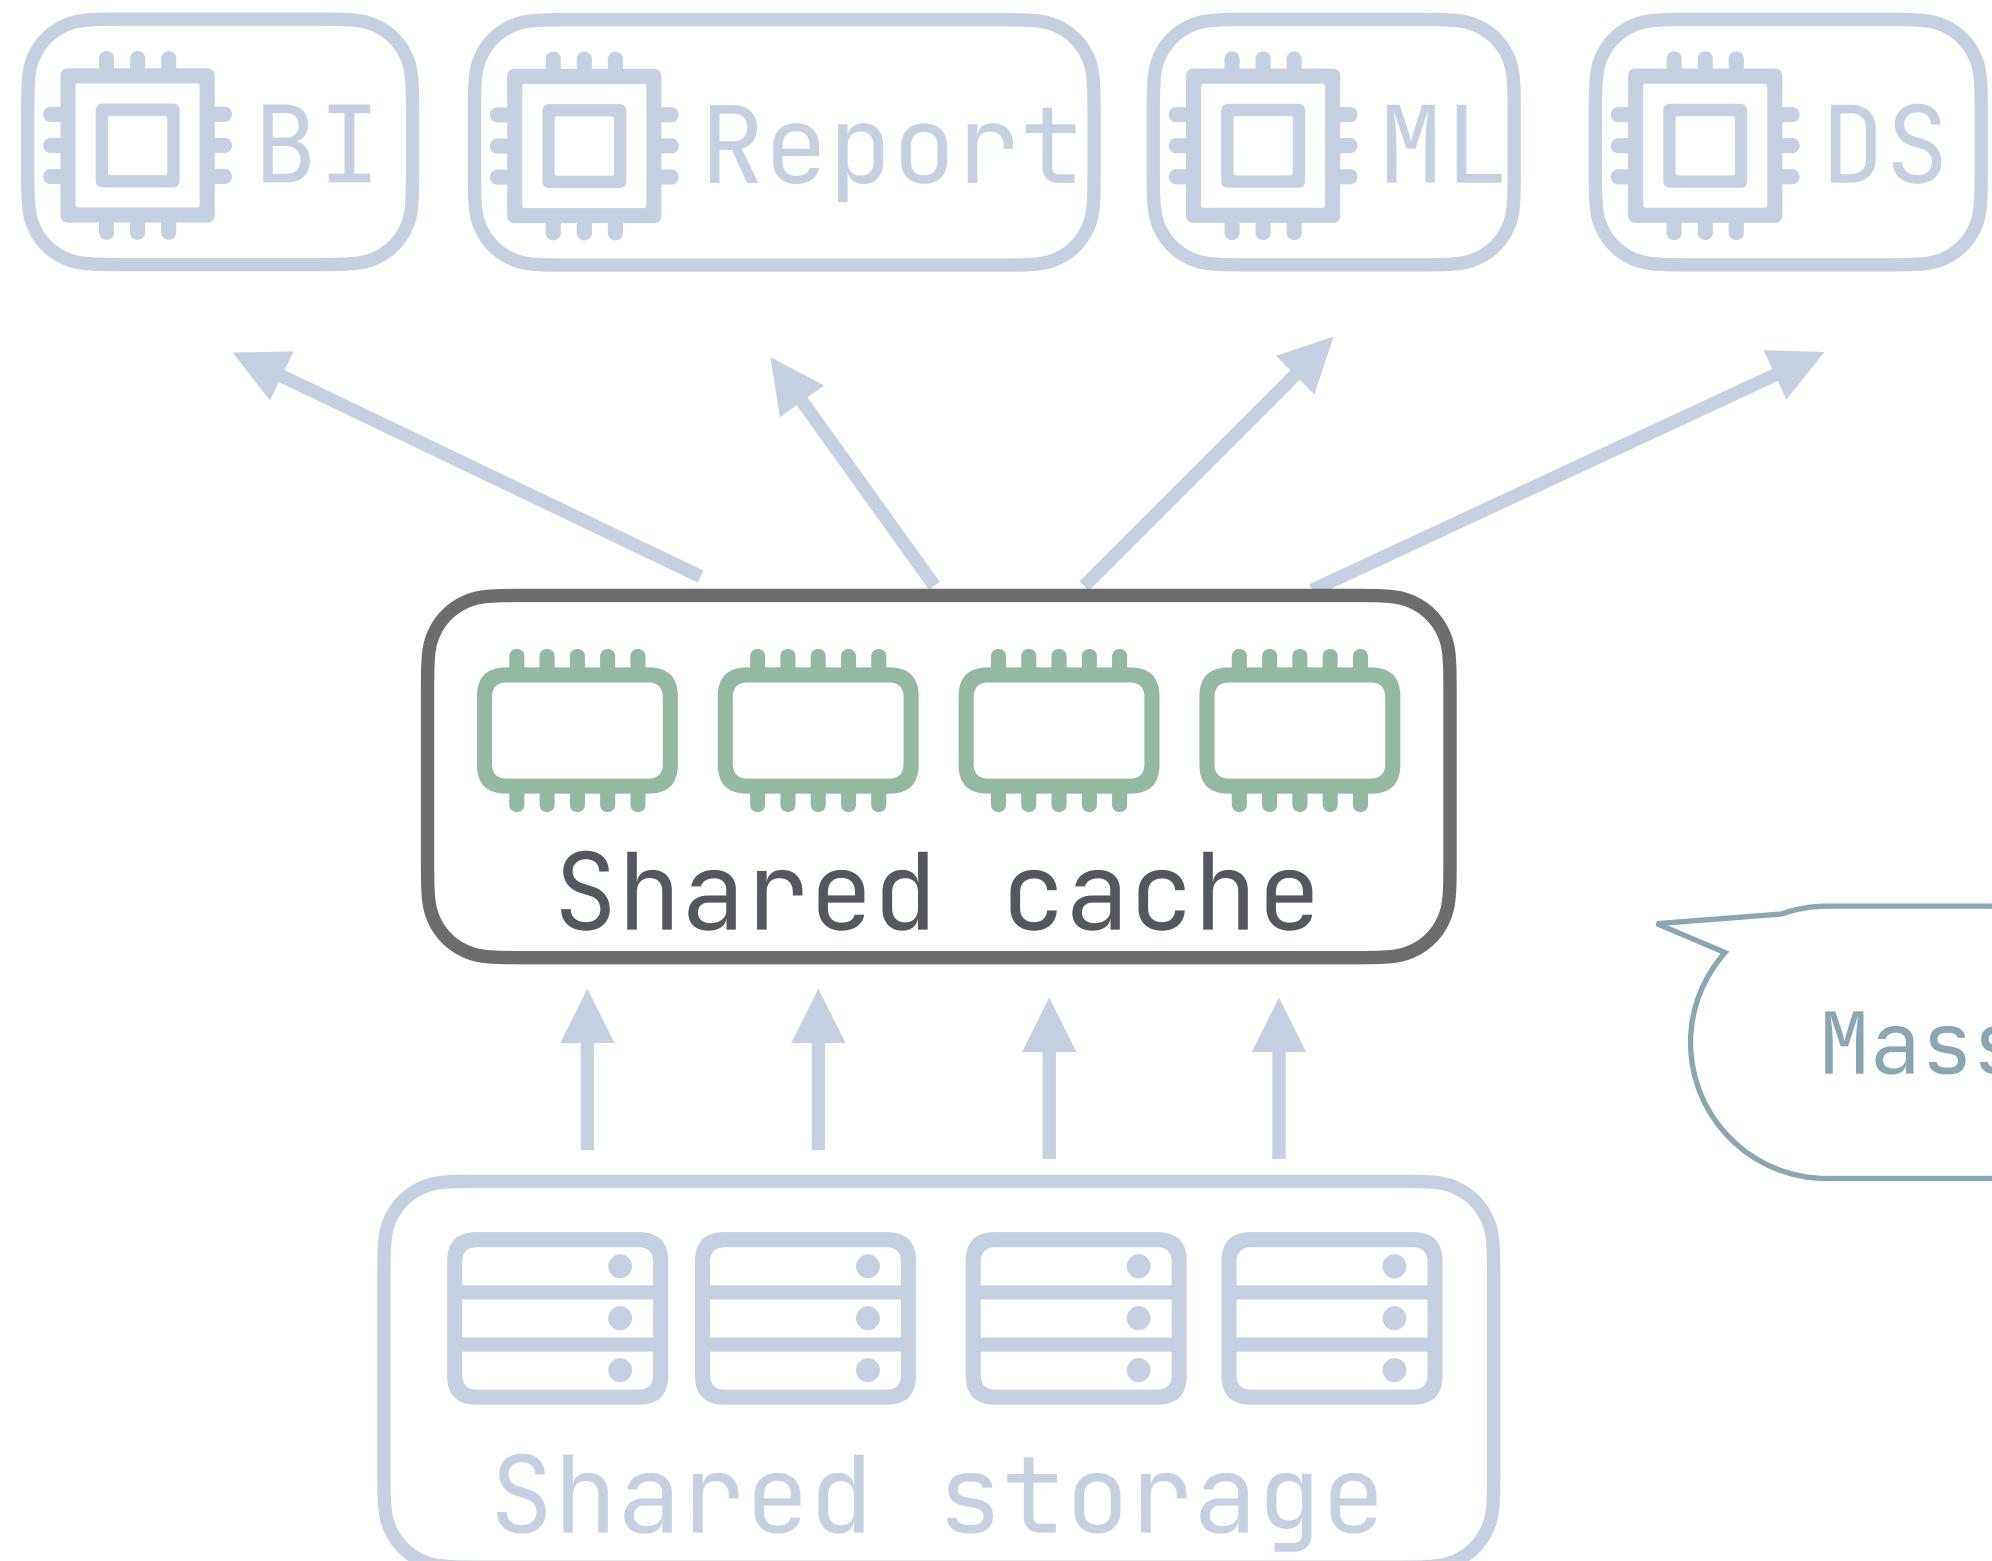
Every system has a cache (2020-2025)



Vision: shared cache (2025+)



First attempt: byte cache

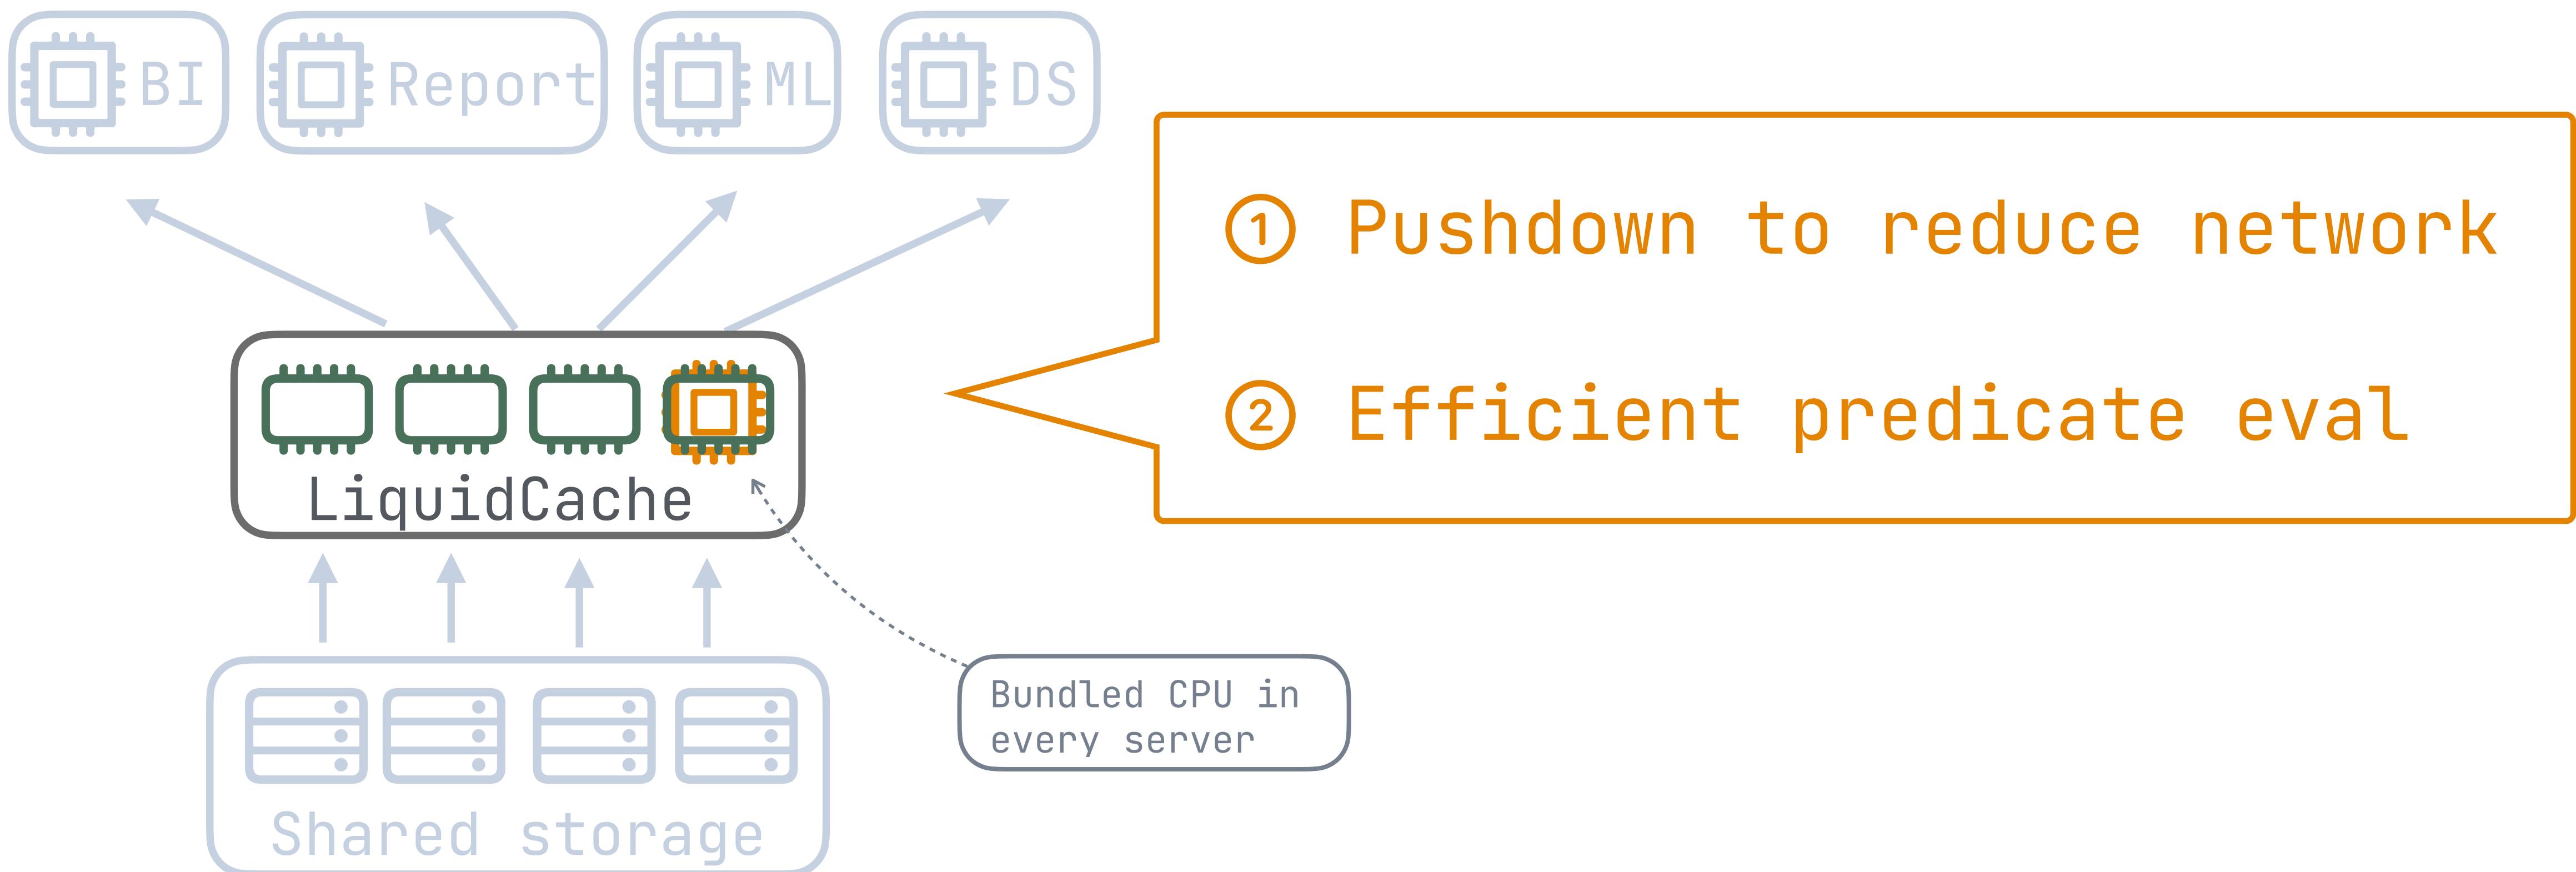


Challenge:
Network bandwidth

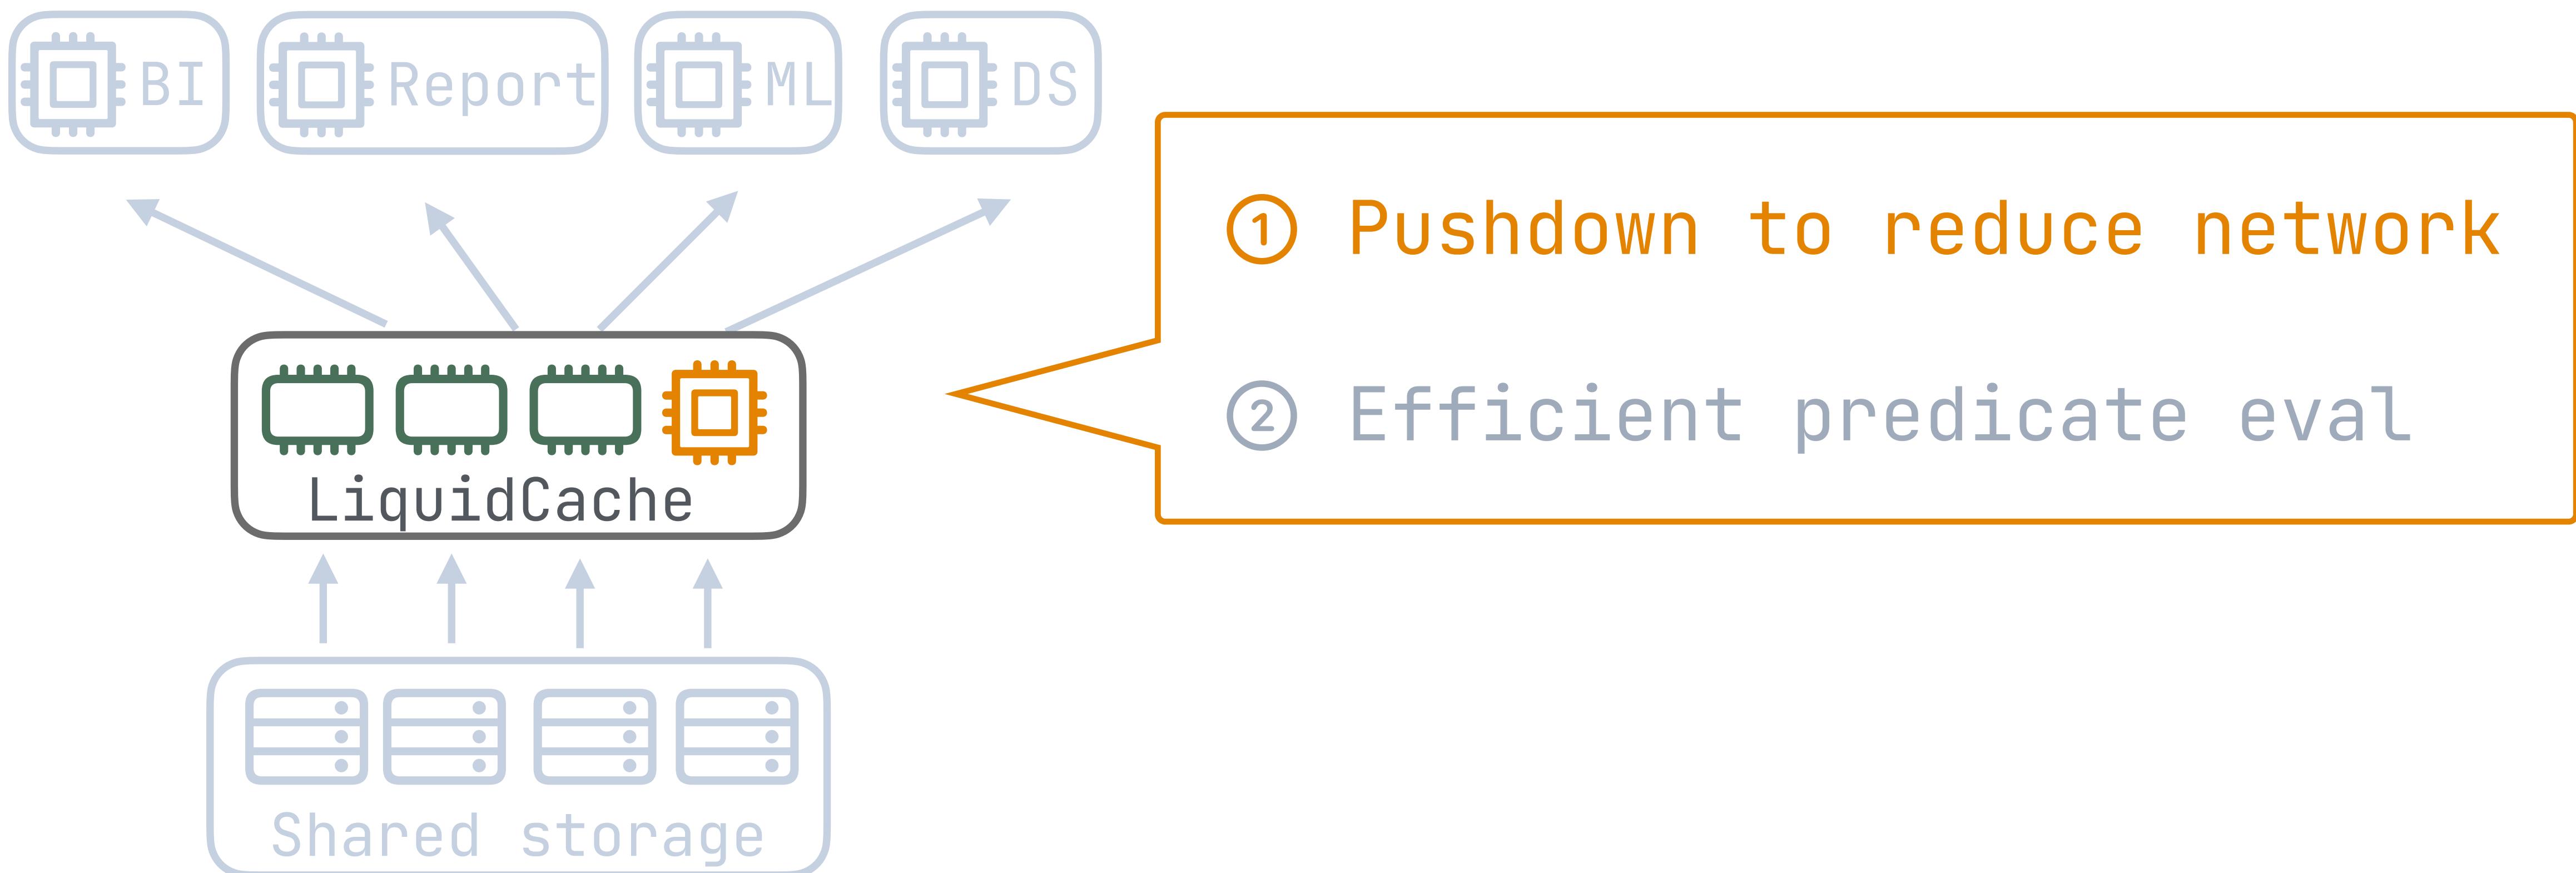
Massive data transfer!

Existing systems: only used
as metadata cache

LiquidCache = compute + data



LiquidCache = compute + data



Prior art: push down

- 1. Access Path Selection in a Relational Database Management System (1979)
- 2. Support for Repetitive Transactions and Ad Hoc Queries in System R (1981)
- 3. Column-Oriented Database System (2009)

Pushdown within query engine

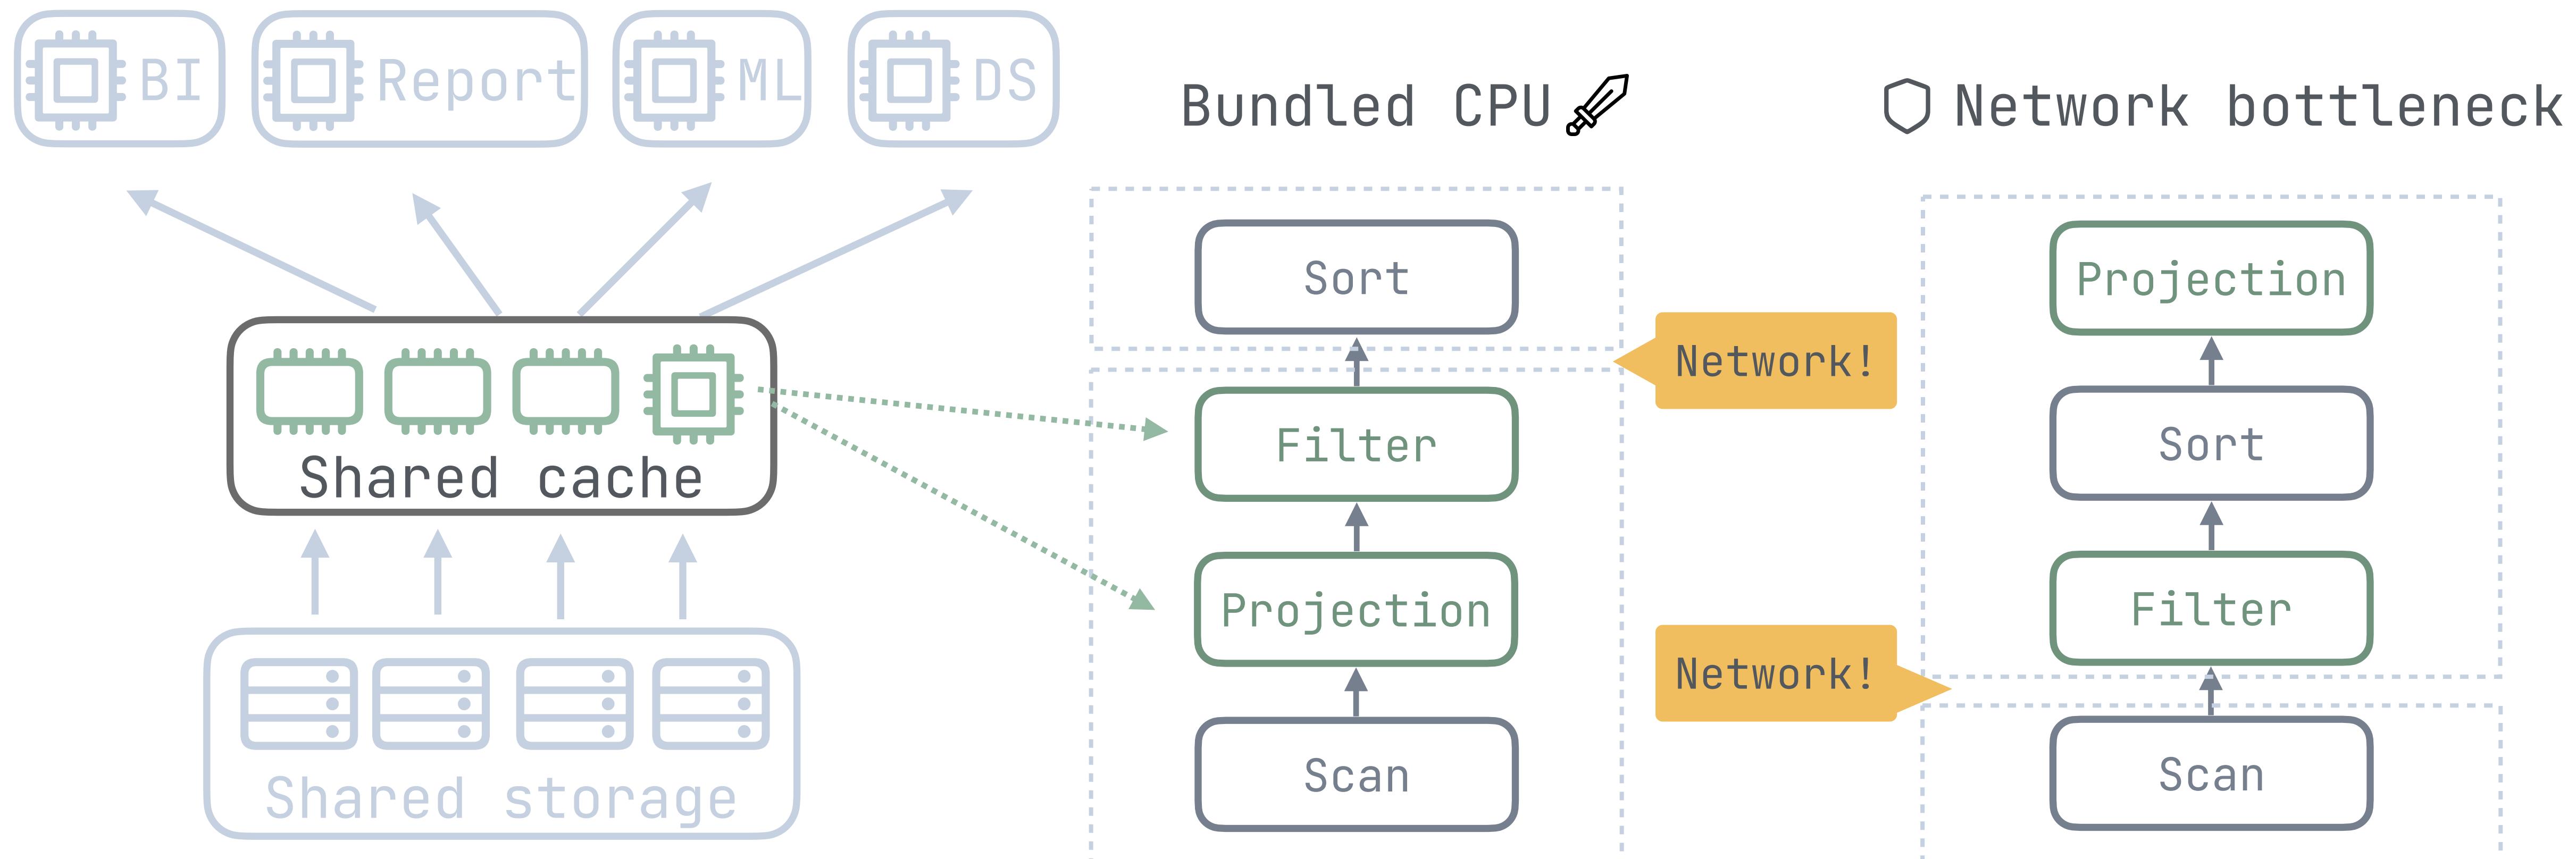
Pushdown to
{NIC, SSD, FPGA, DPU}

- 1. Query Processing on Smart SSDs: Opportunities and Challenges (2013)
- 2. Introduction to the IBM Netezza warehouse appliance. (2011)
- 3. DPDPUs: Data Processing with DPUs. (2024)
- 4. Towards Accelerating Data Intensive Application's Shuffle Process Using SmartNICs (2023)

- 1. PushdownDB: Accelerating a DBMS using S3 Computation (2020)
- 2. FlexPushdownDB: Hybrid Pushdown and Caching in a Cloud DBMS (2021)
- 3. Fusion: An Analytics Object Store Optimized for Query Pushdown. (2025)

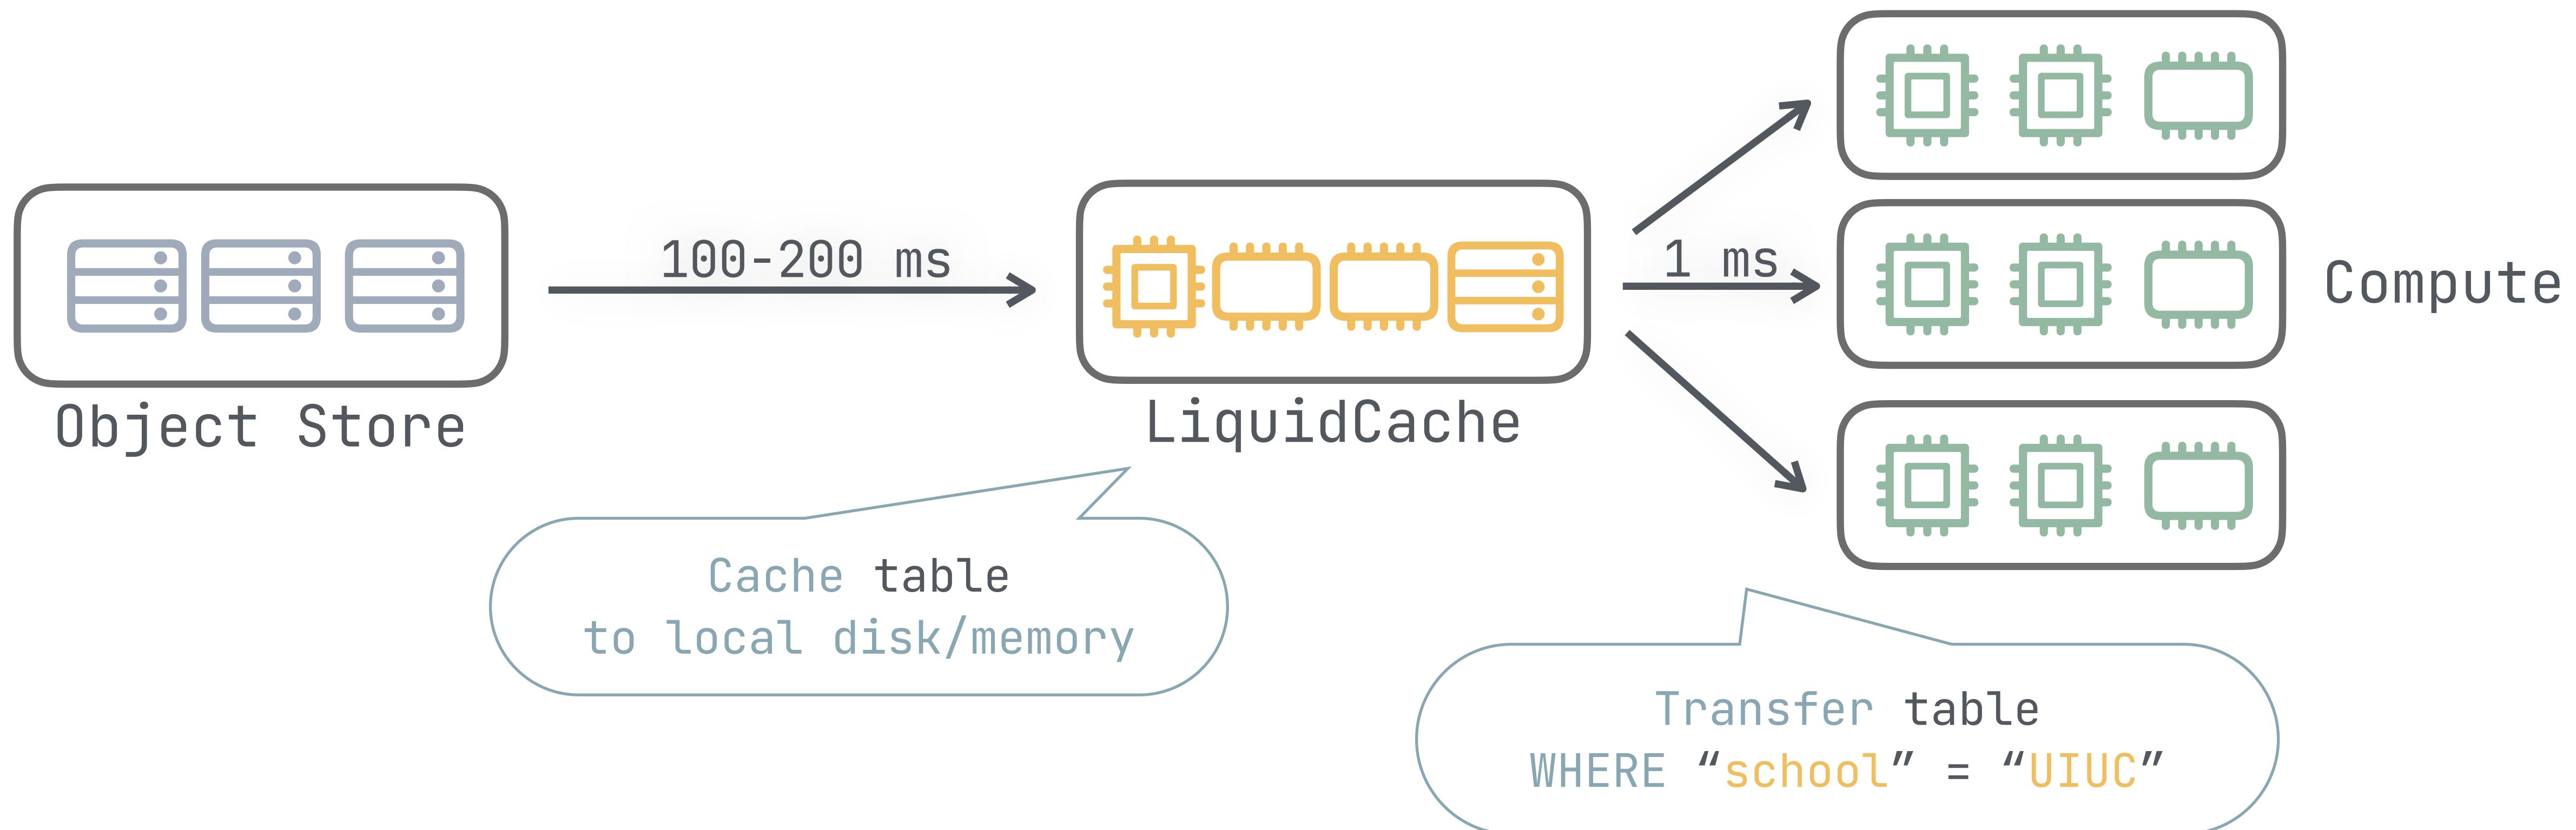
Pushdown to
Cloud storage

Pushdown to cache service

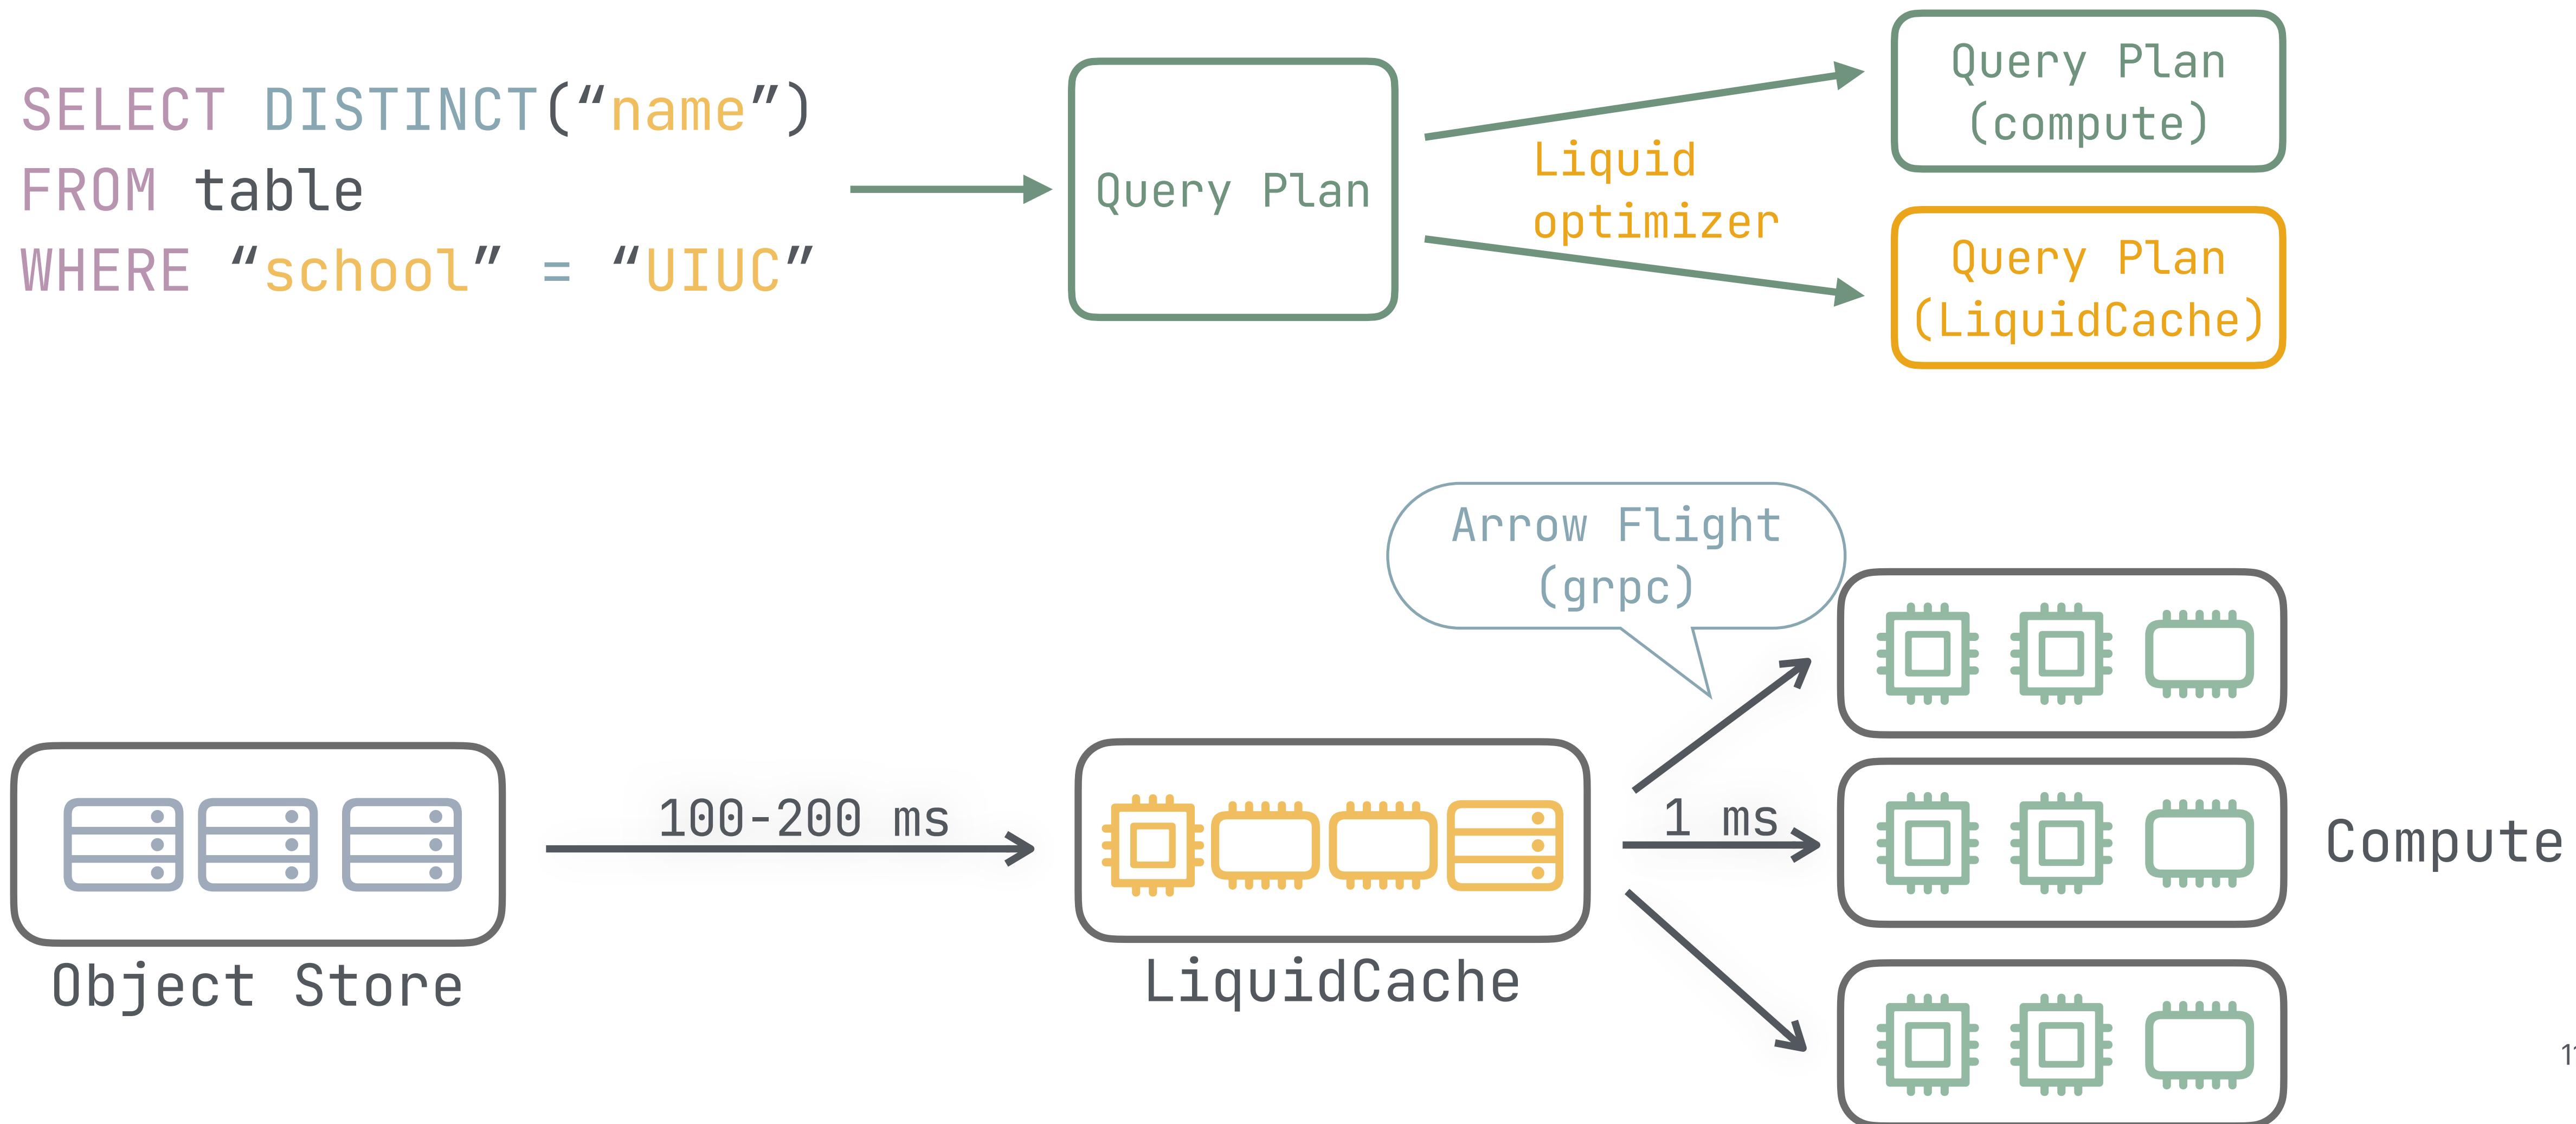


Pushdown example

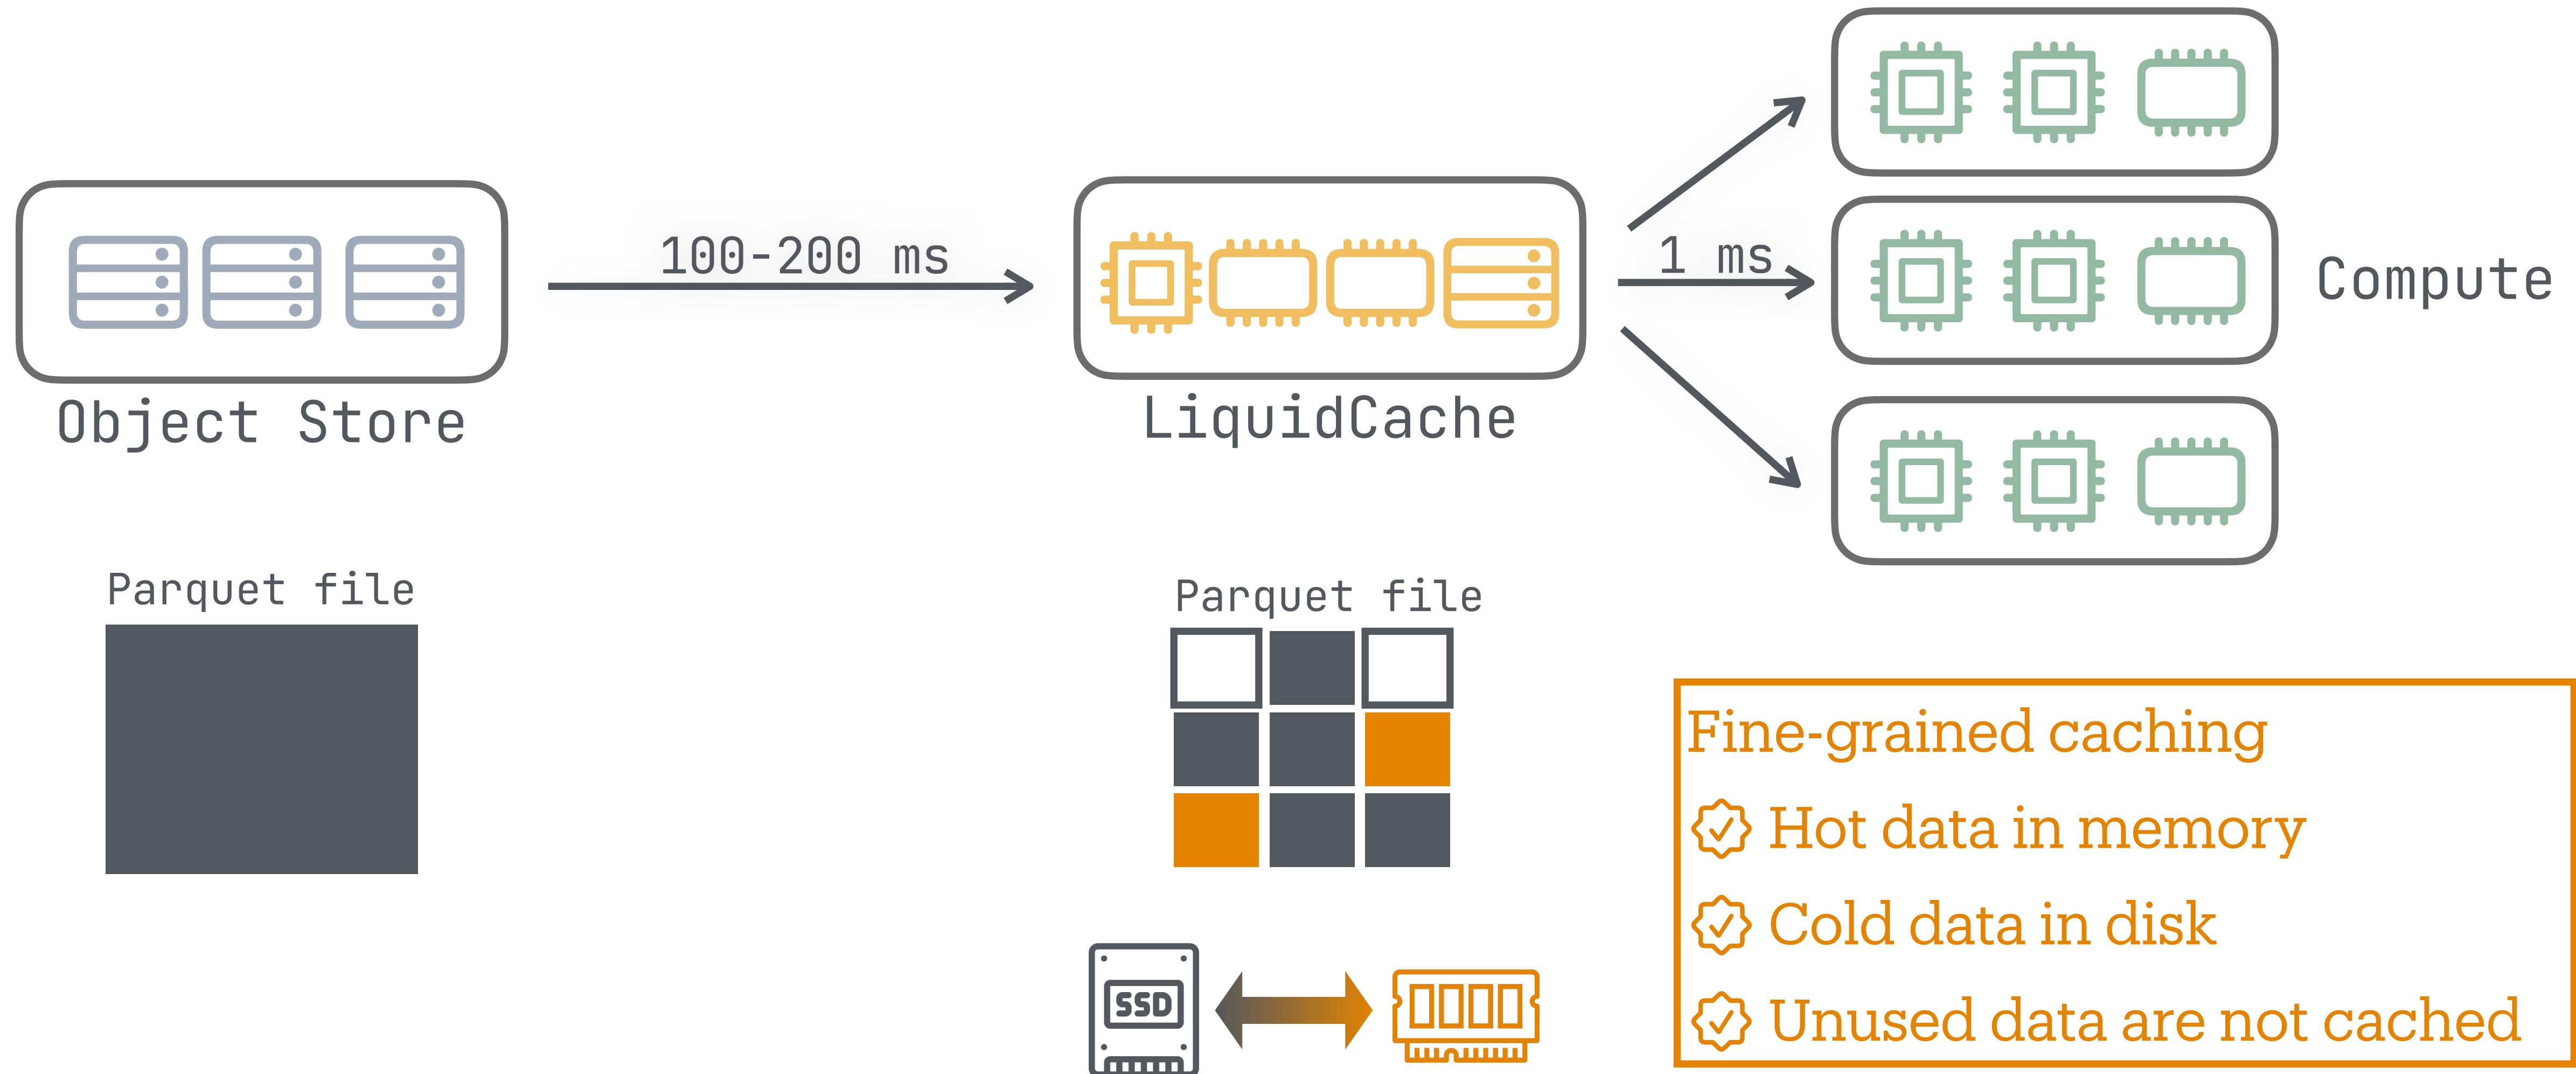
```
SELECT DISTINCT("name")
FROM table
WHERE "school" = "UIUC"
```



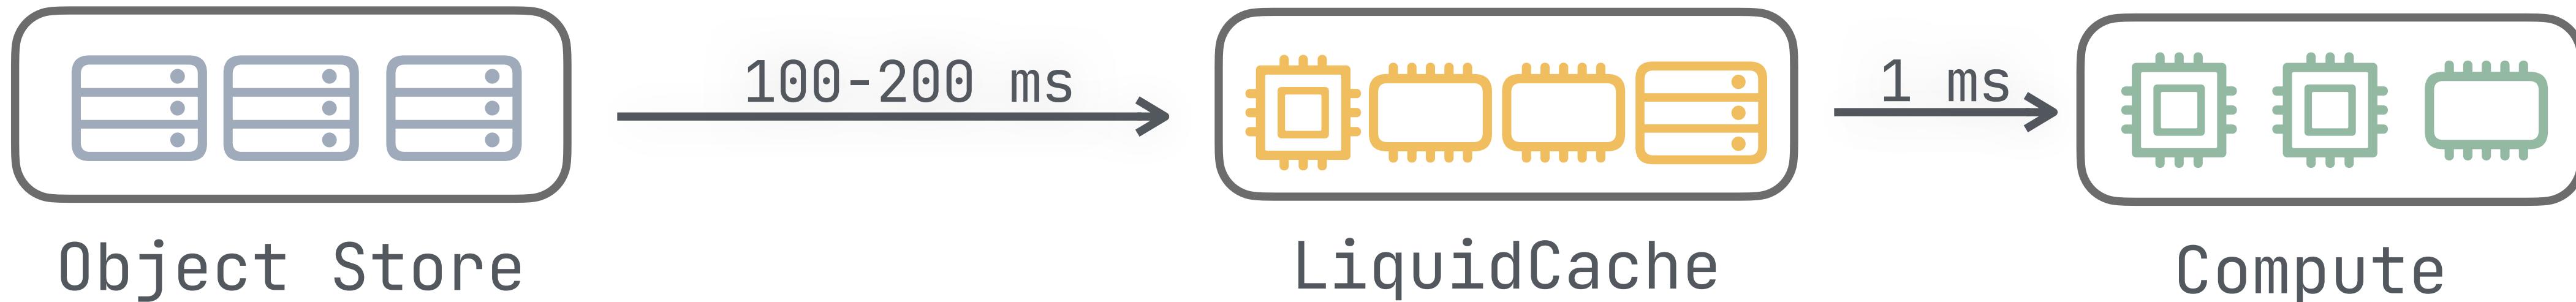
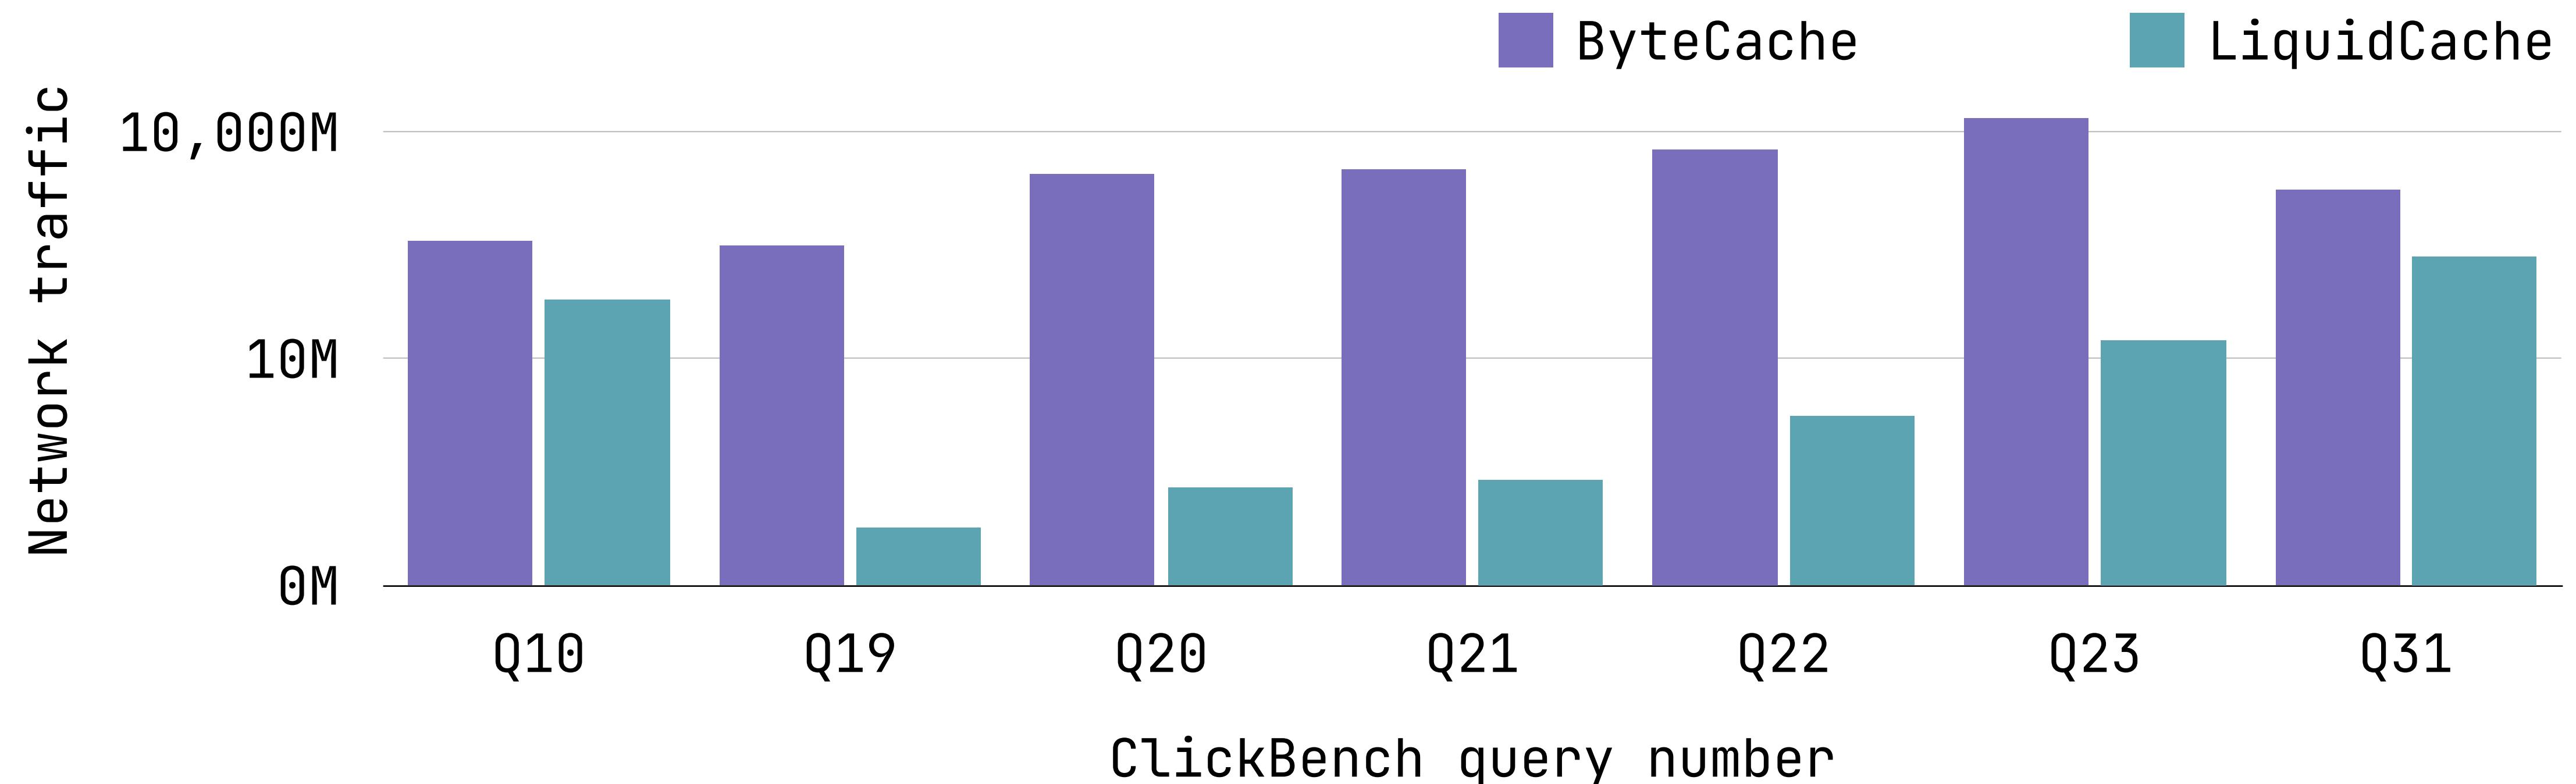
LiquidCache Architecture



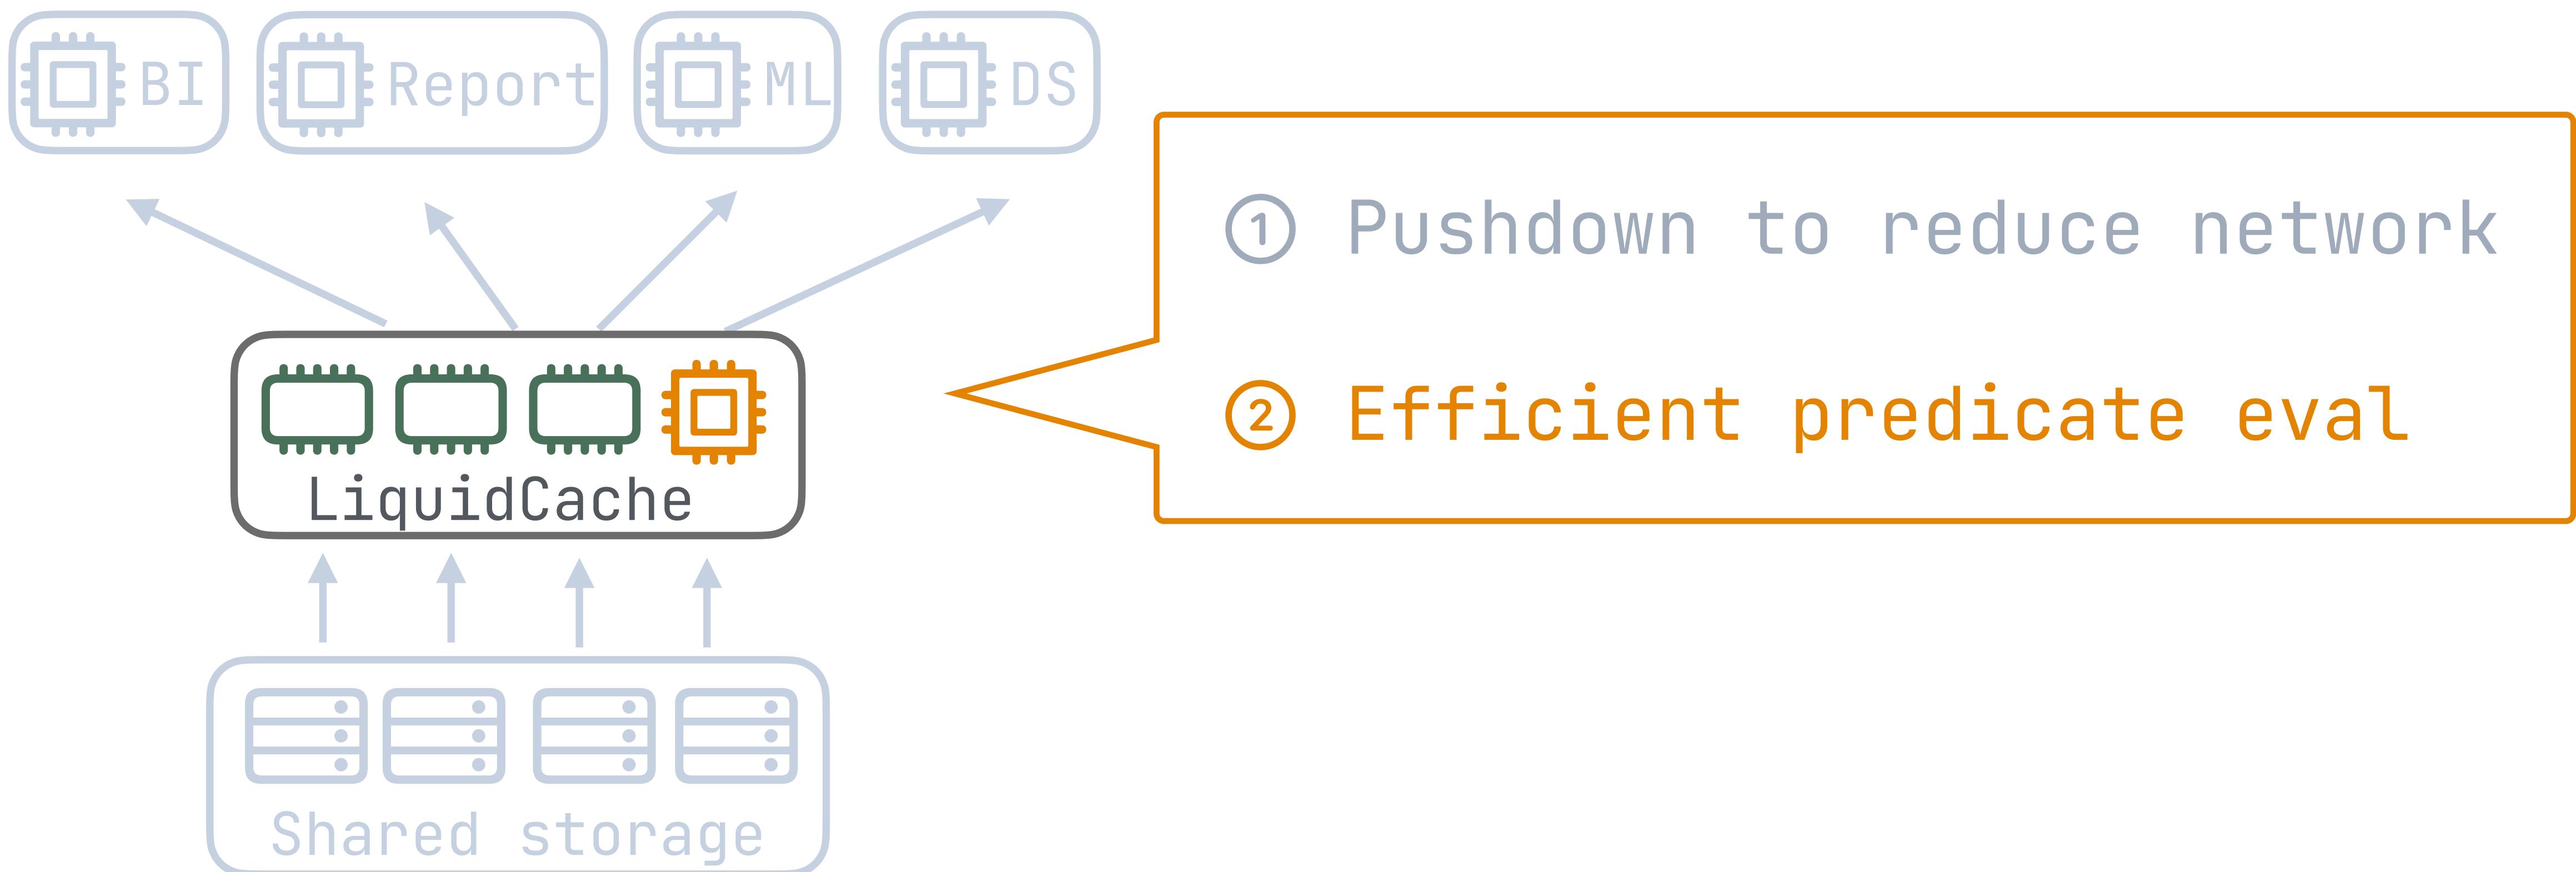
Fine-grained caching



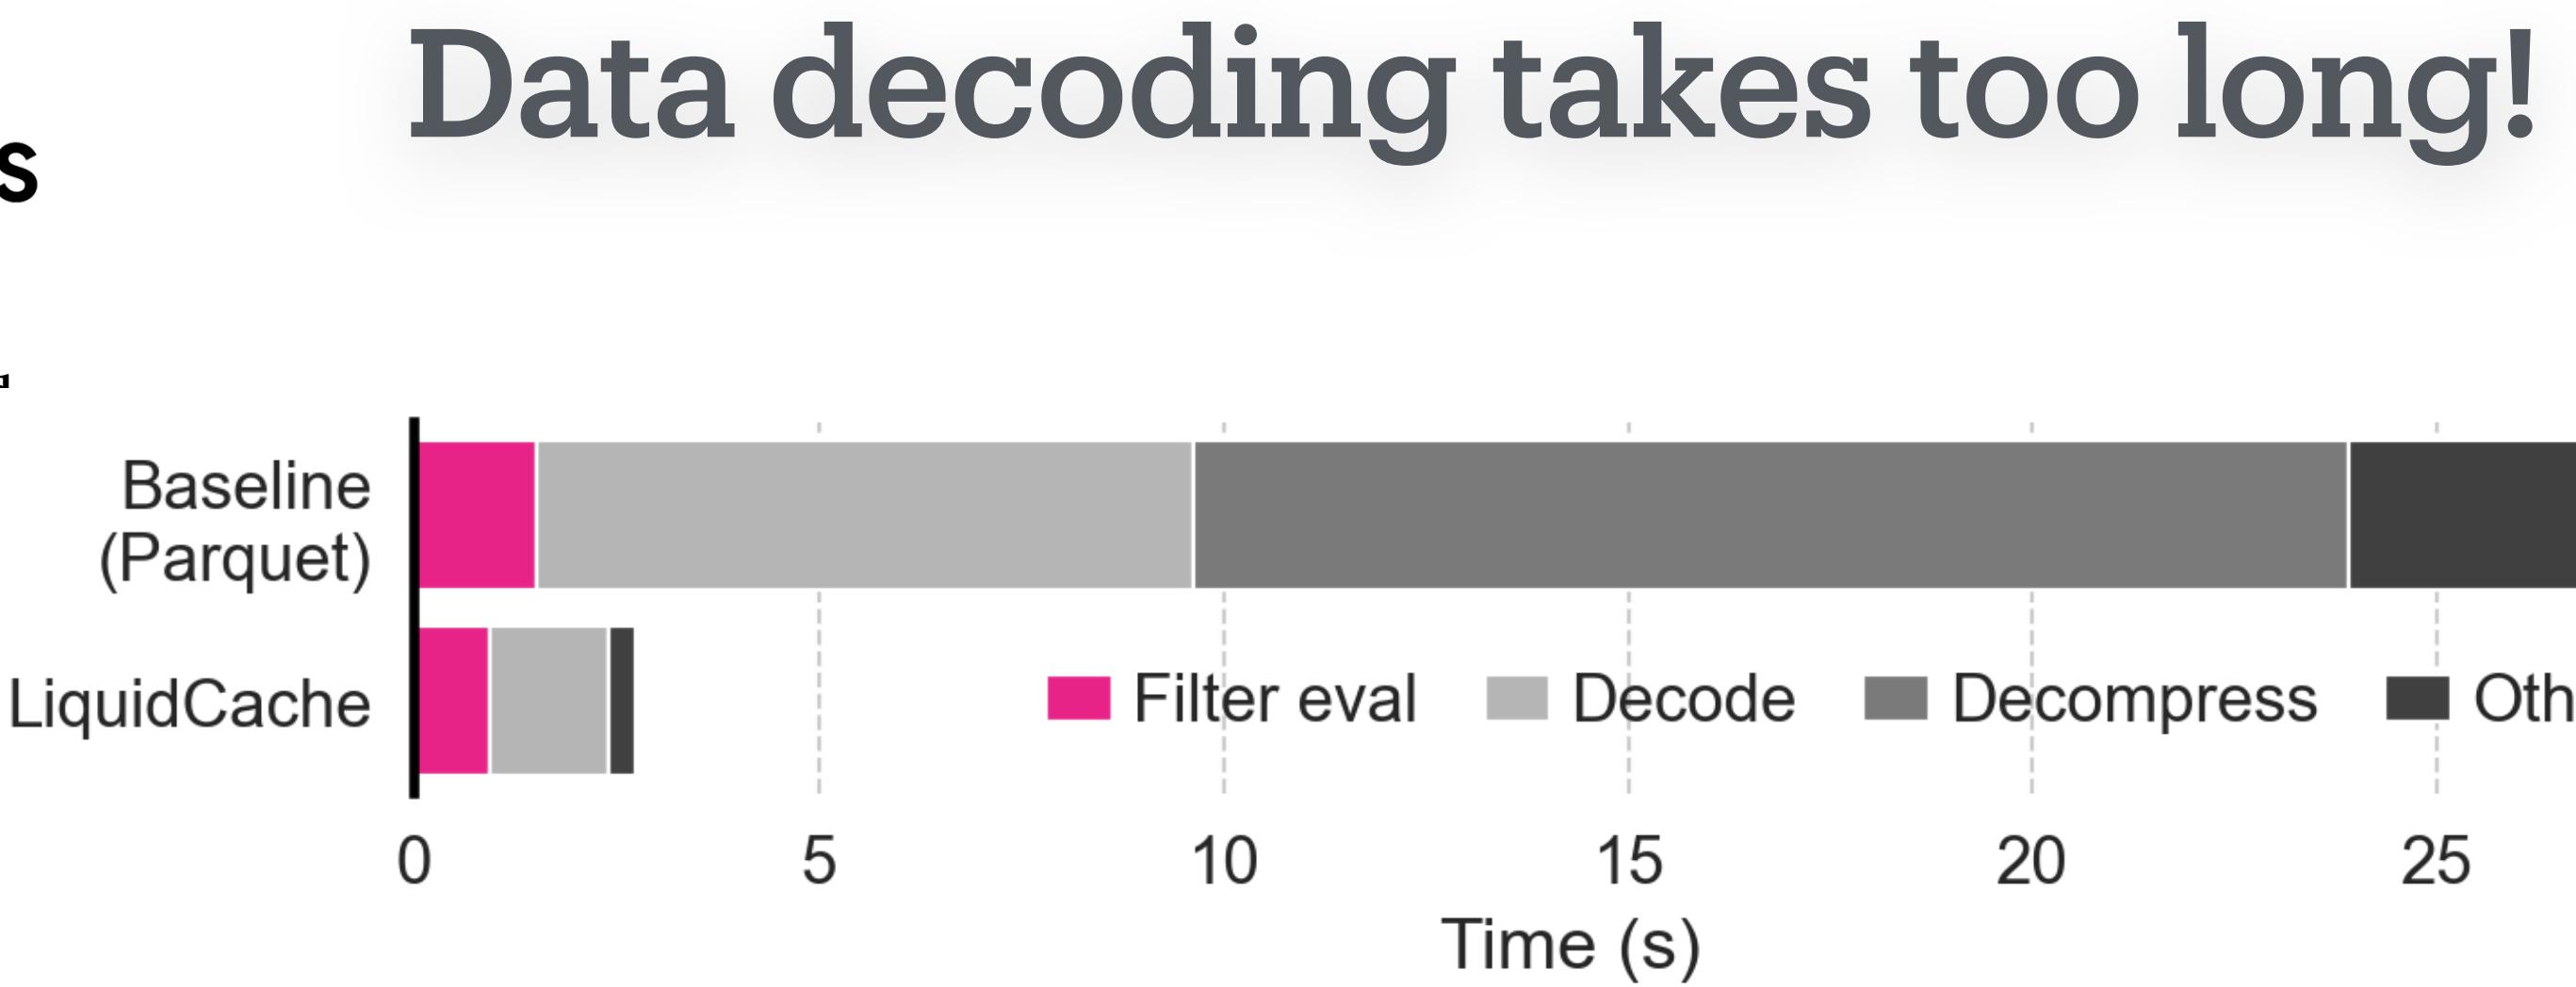
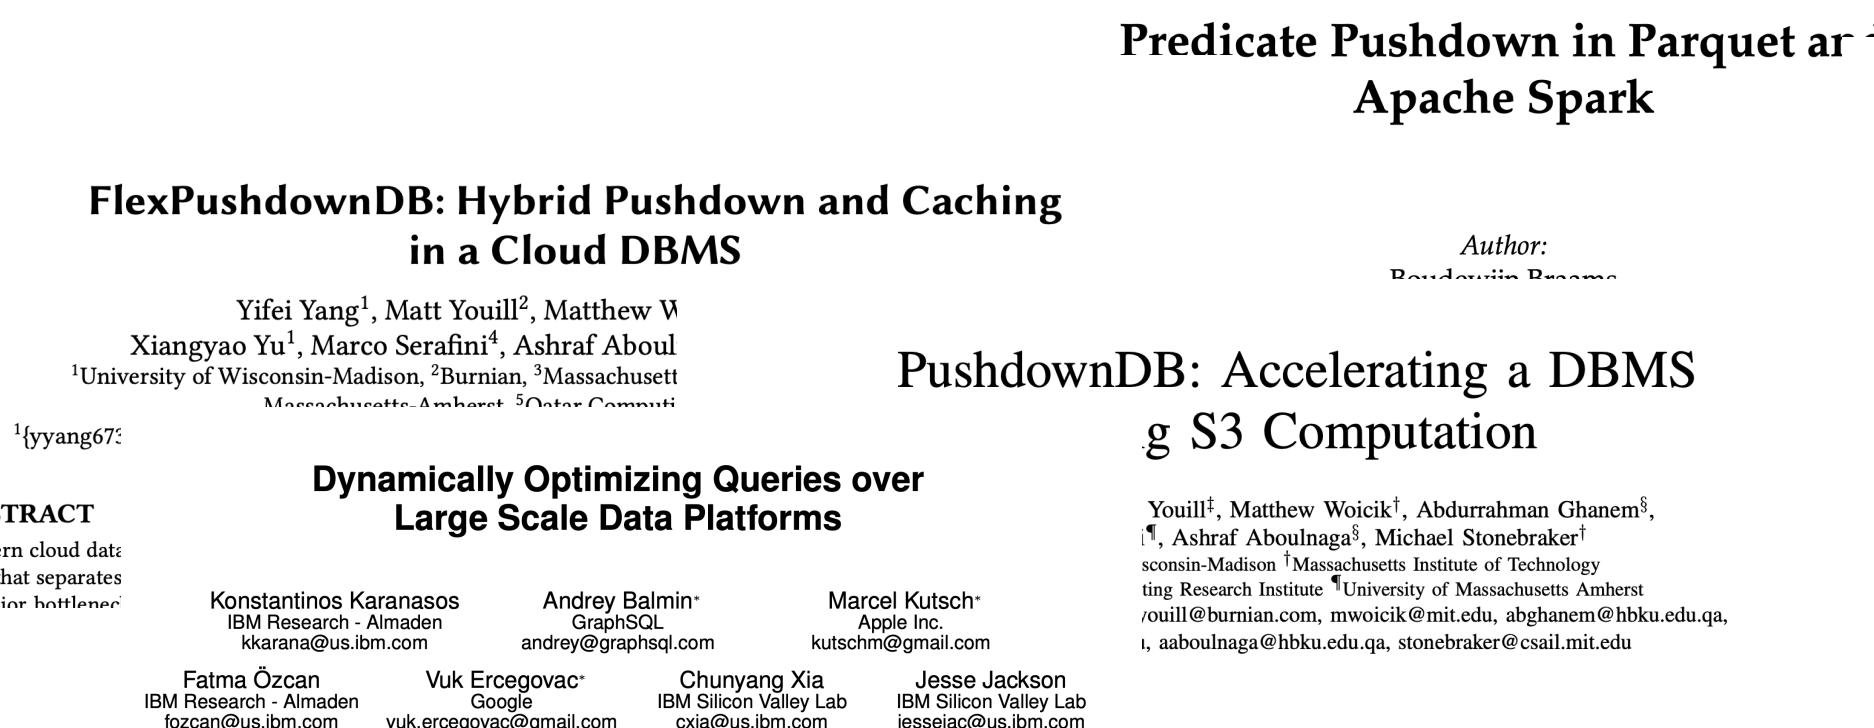
Push down reduce network traffic



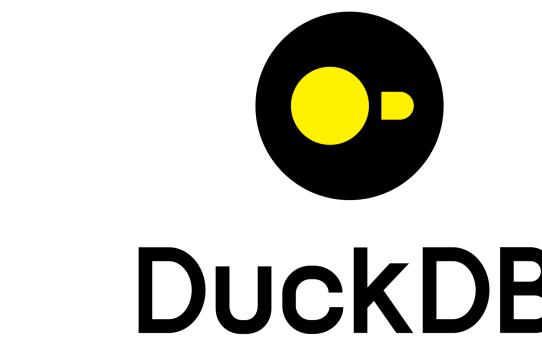
LiquidCache = compute + data



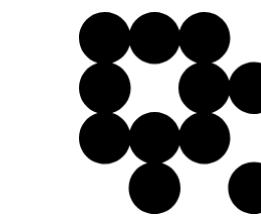
Pushdown is expensive



Prior Art: “better” file formats



DuckDB

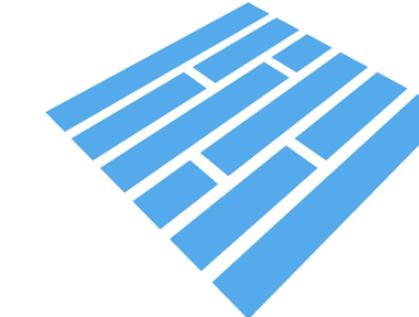


LanceDB



Vortex

ClickHouse



Parquet 2.0

Industry makes
new file formats

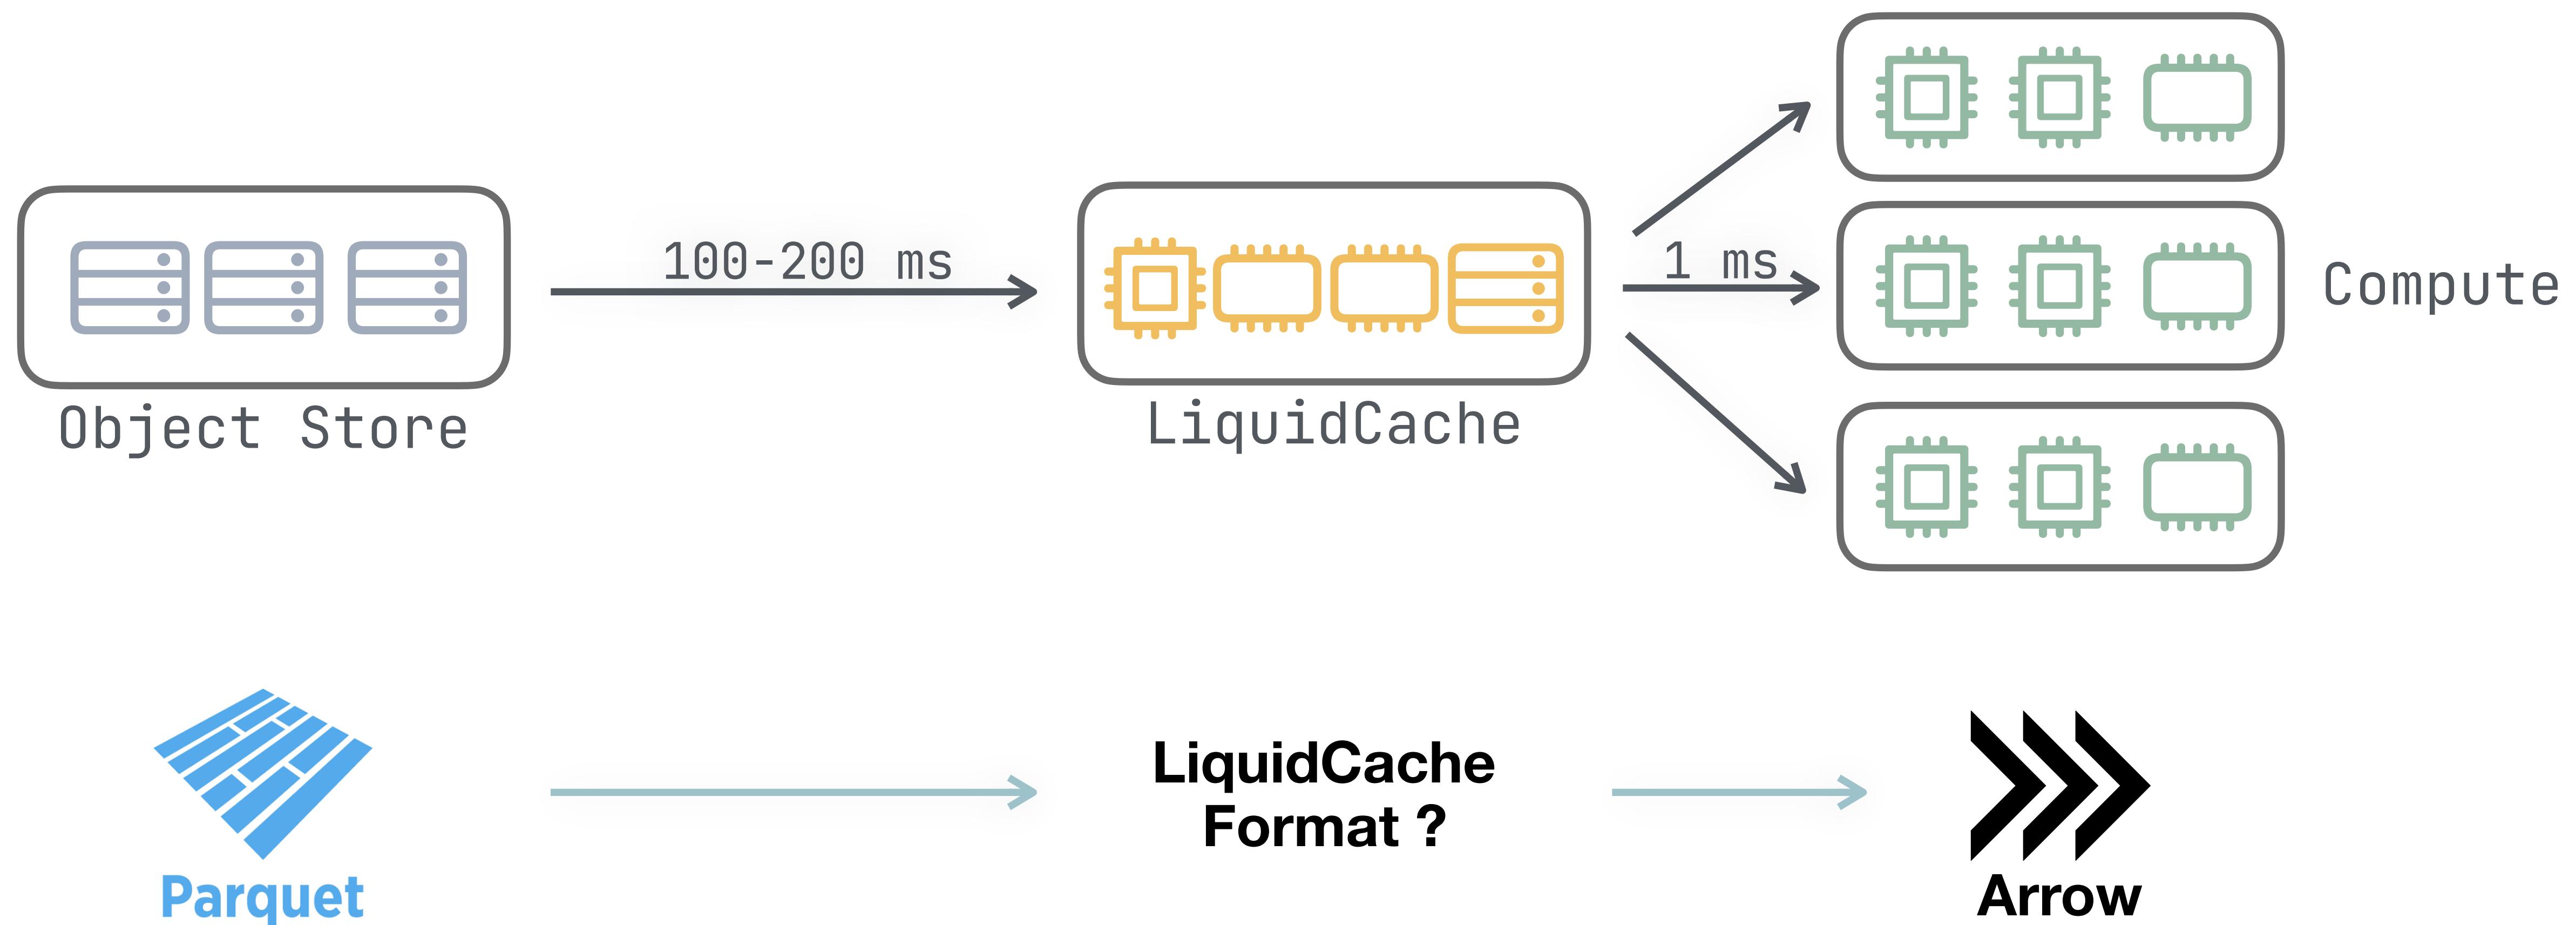
1. FSST: Fast Random Access String Compression (2020)
2. BtrBlocks: Efficient Columnar Compression for Data Lakes (2023)
3. The Fast Lanes Compression Layout: Decoding 100 Billion Integers per Second with Scalar Code (2023)
4. ALP: Adaptive Lossless floating-Point Compression (2023)
5. The FastLanes File Format (2025)
6. AnyBlox: A Framework for Self-Decoding Datasets (2025)
7. F3: The Open-Source Data File Format for the Future (2025)

Researchers propose
new encodings

Battle tested
Open & stable governance
Cross-engine sharing
Reasonable performance

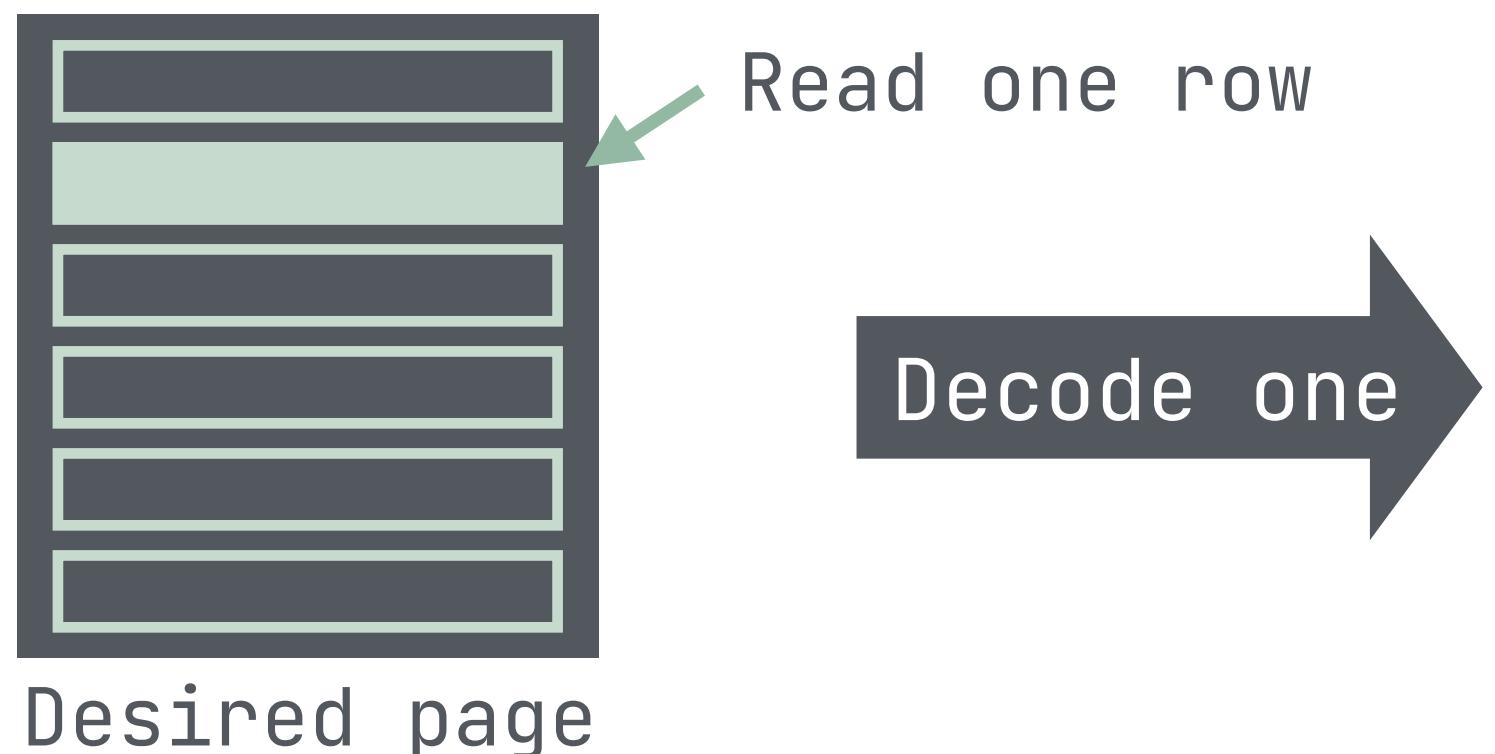
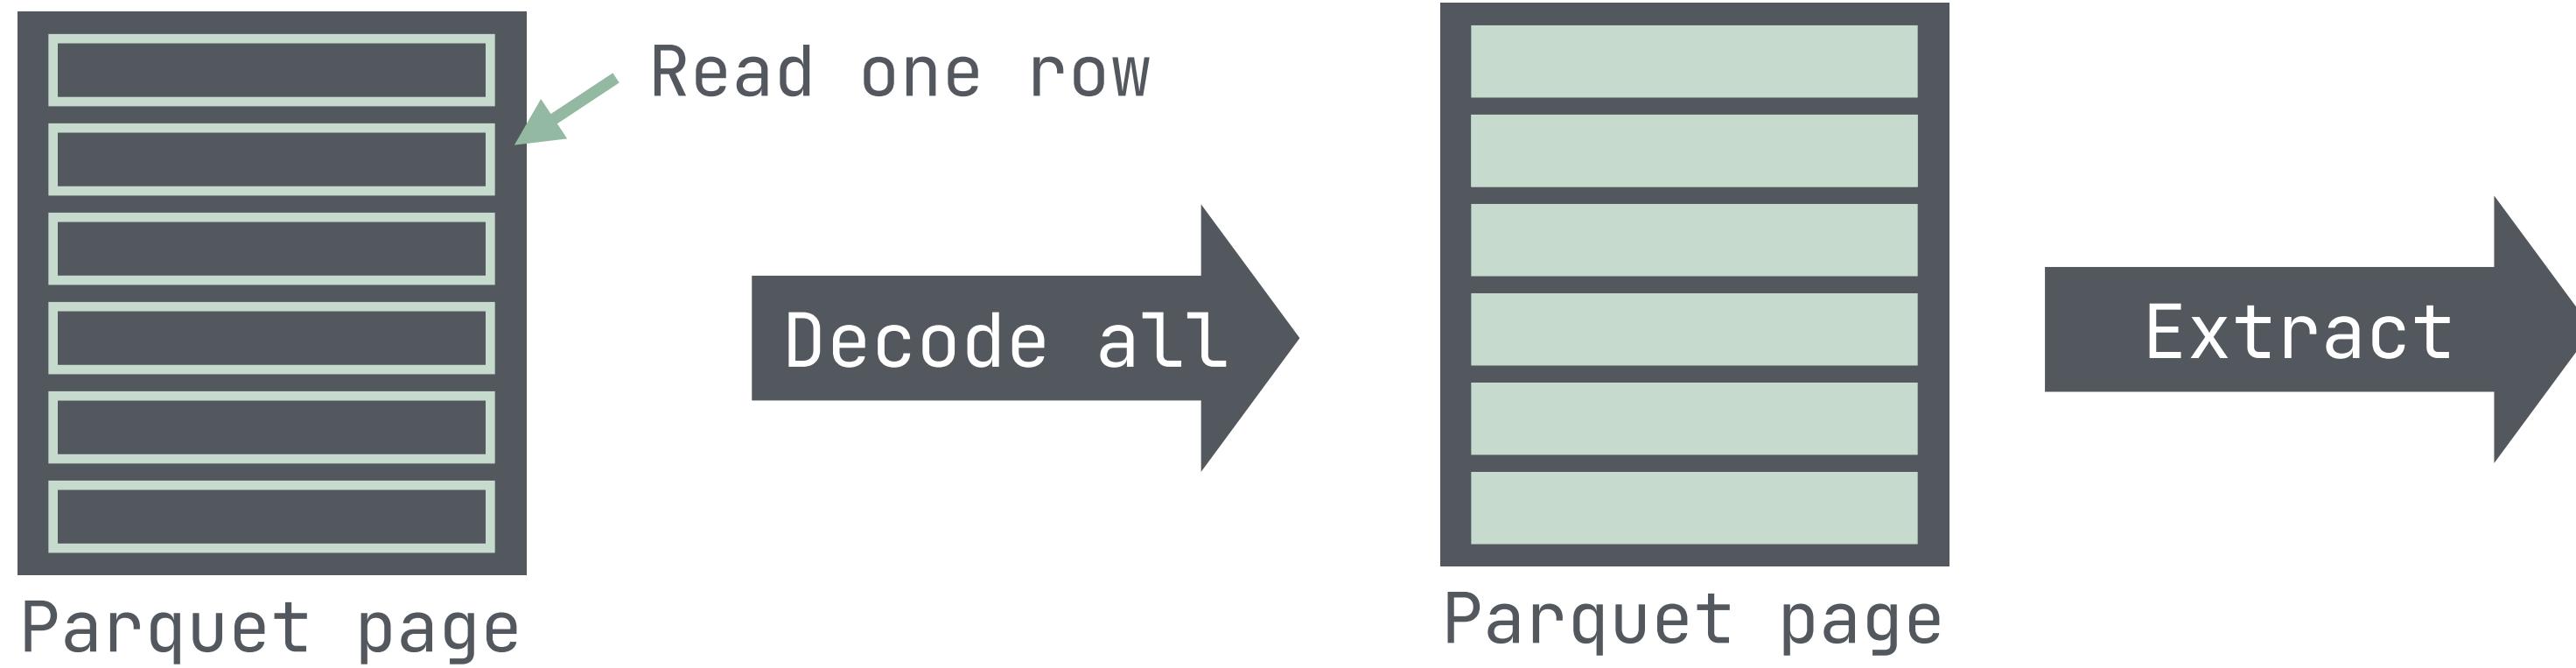
But Industry is
locked to old Parquet

LiquidCache: cache-only format

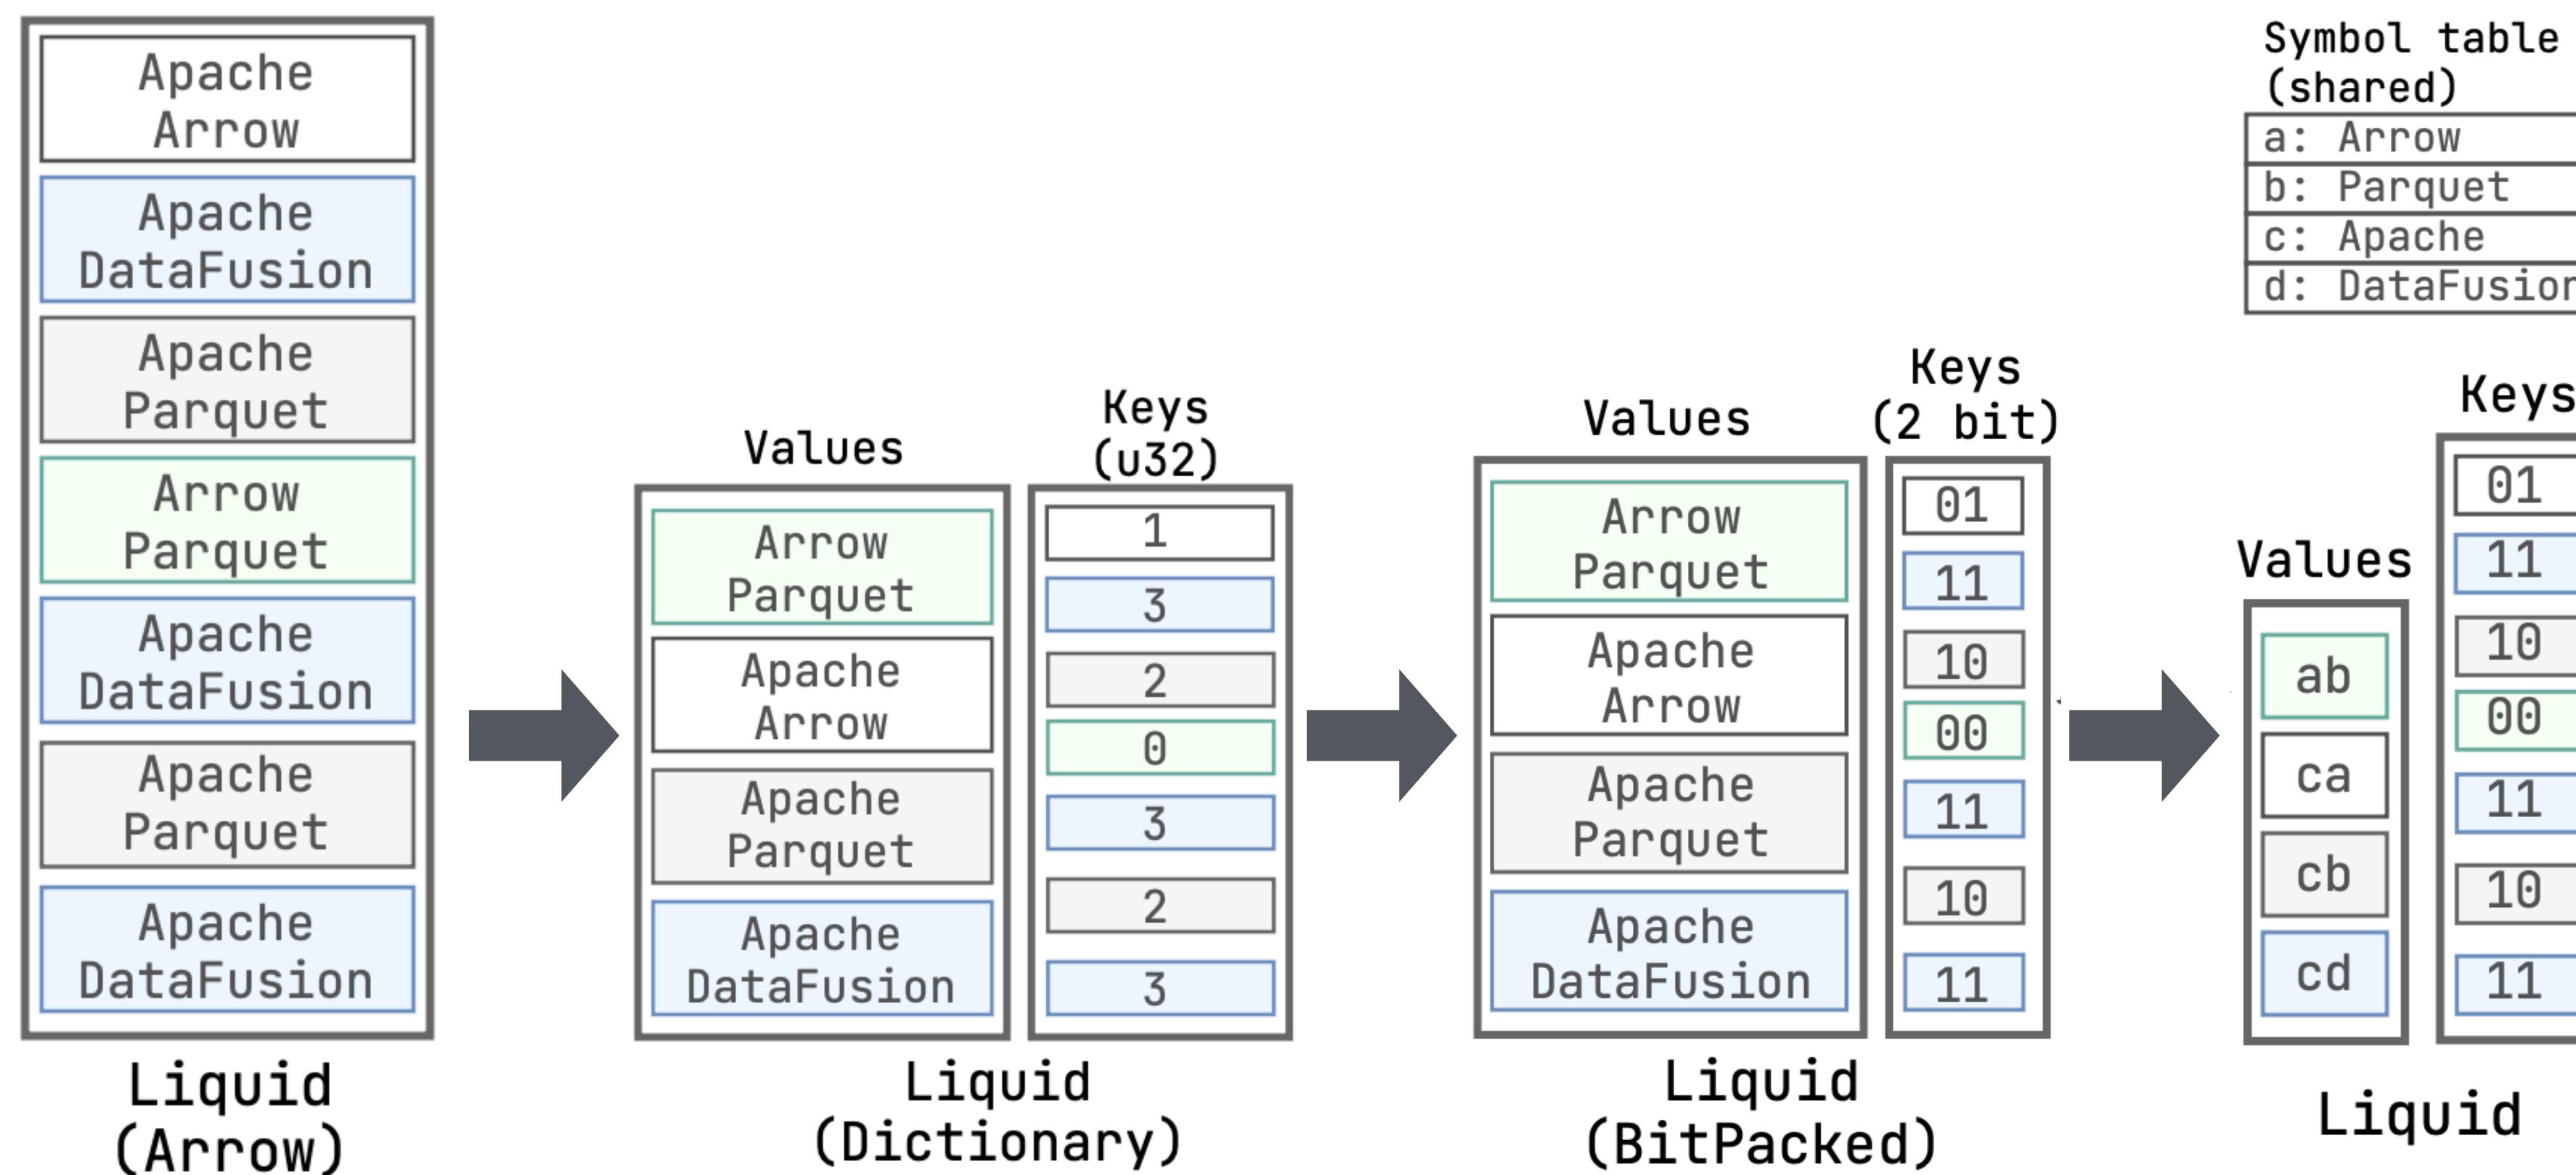


Principle:

Each row must be independently decodable



Each row must be independently decodable (string example)



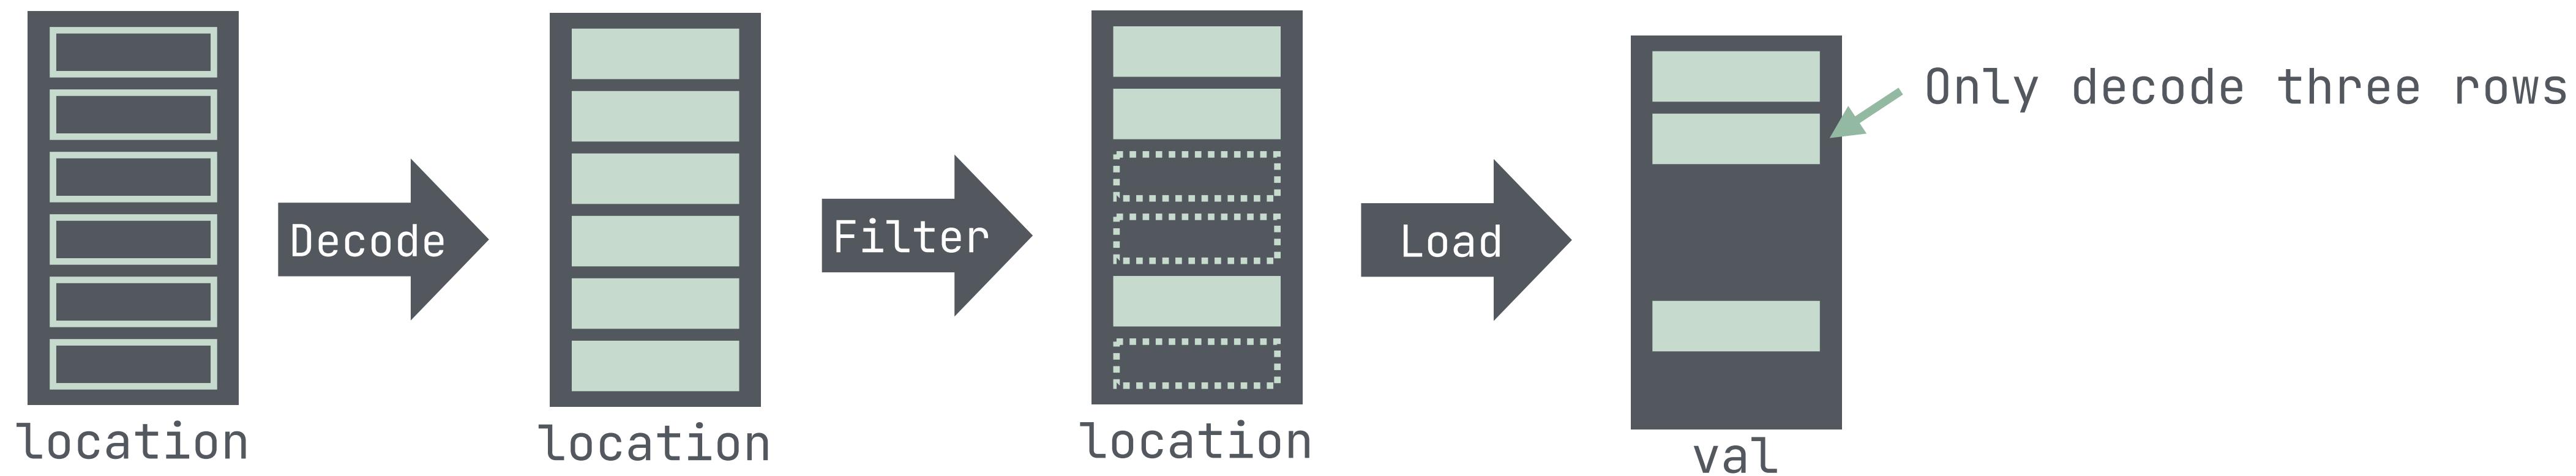
No general purpose compression

Leverages state-of-the-art encoding schemes

Carefully designed encoding/layout for each data types

Co-design with filter pushdown (selective decoding)

```
SELECT val, location  
FROM sensor_data  
WHERE location = 'office';
```



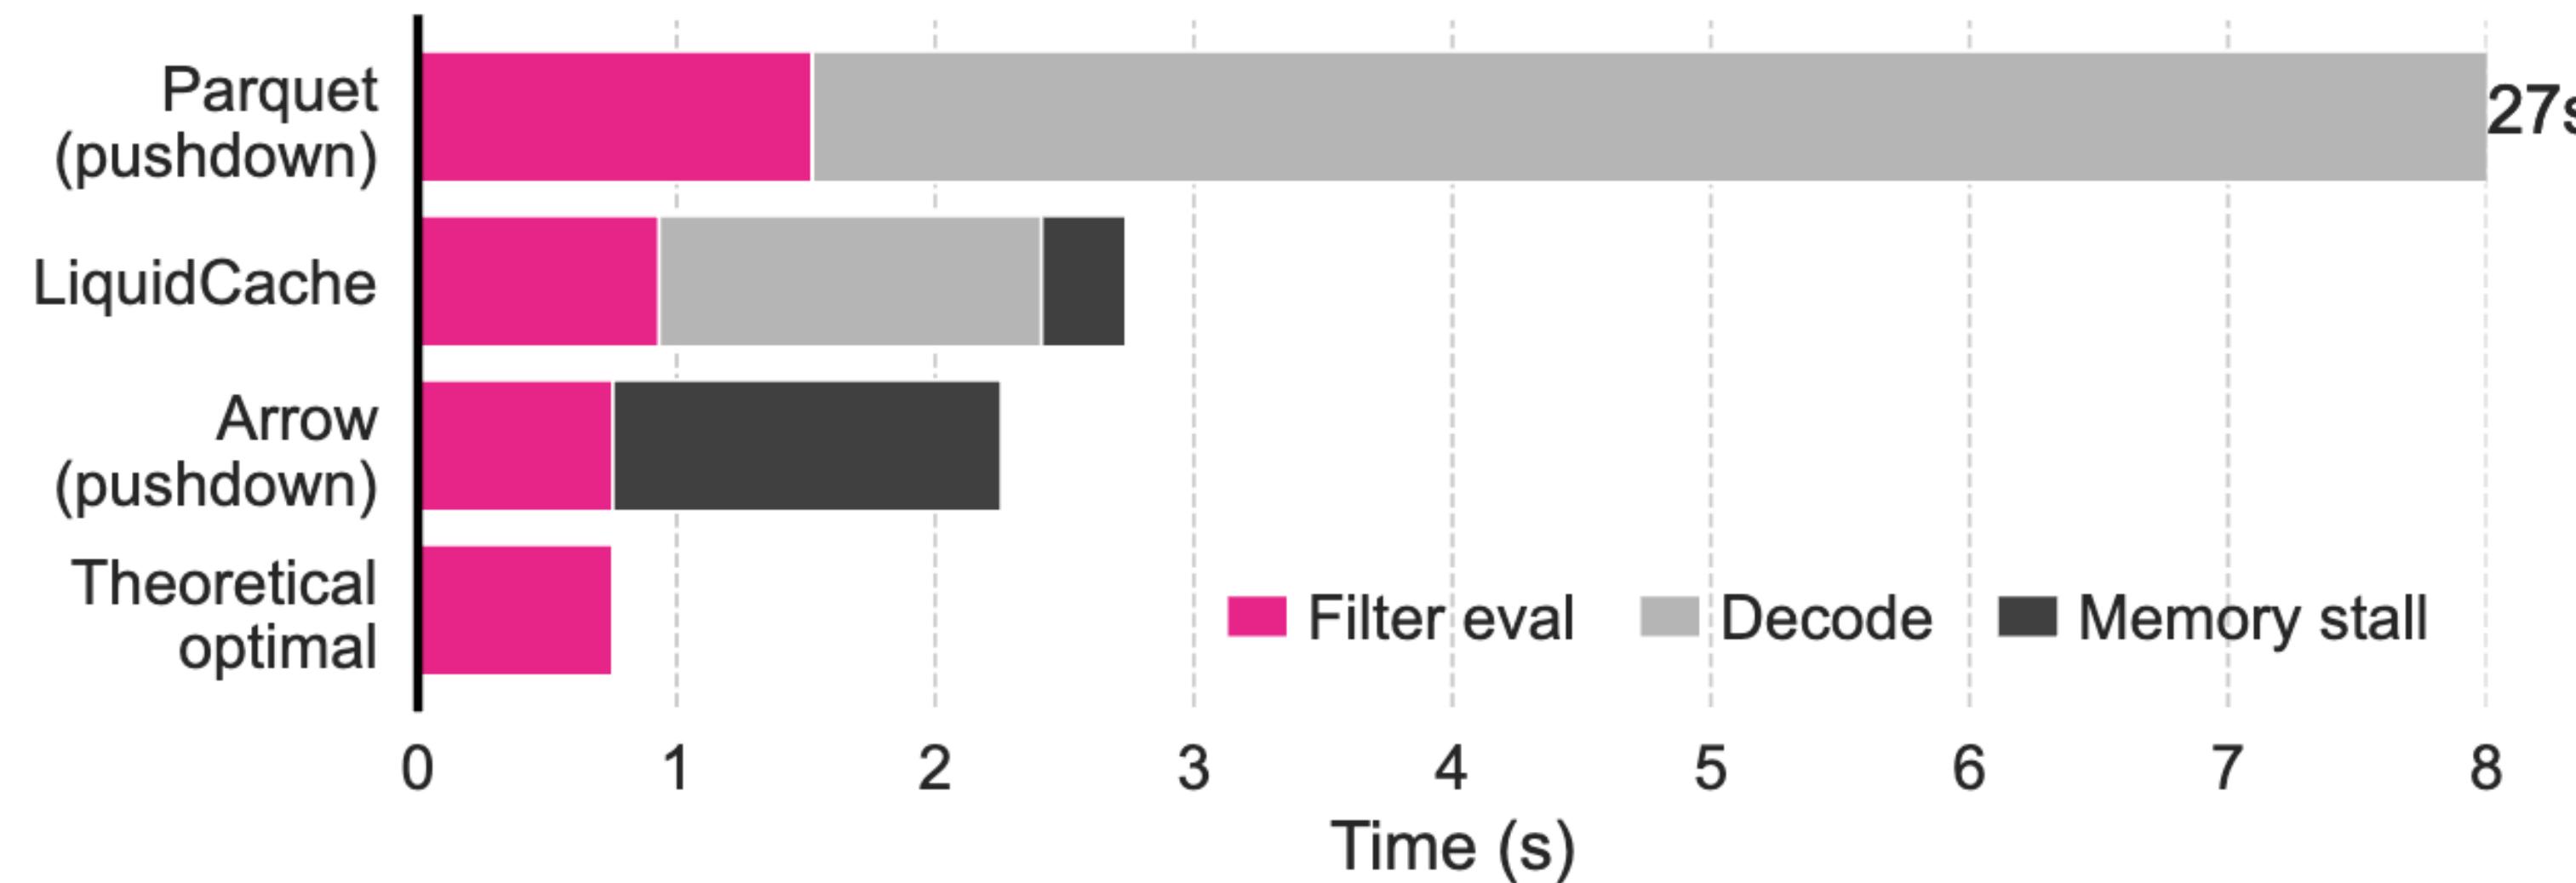
More in our paper:

- Late materialization
- Evaluate on encoded data
- Partial decoding



LiquidCache

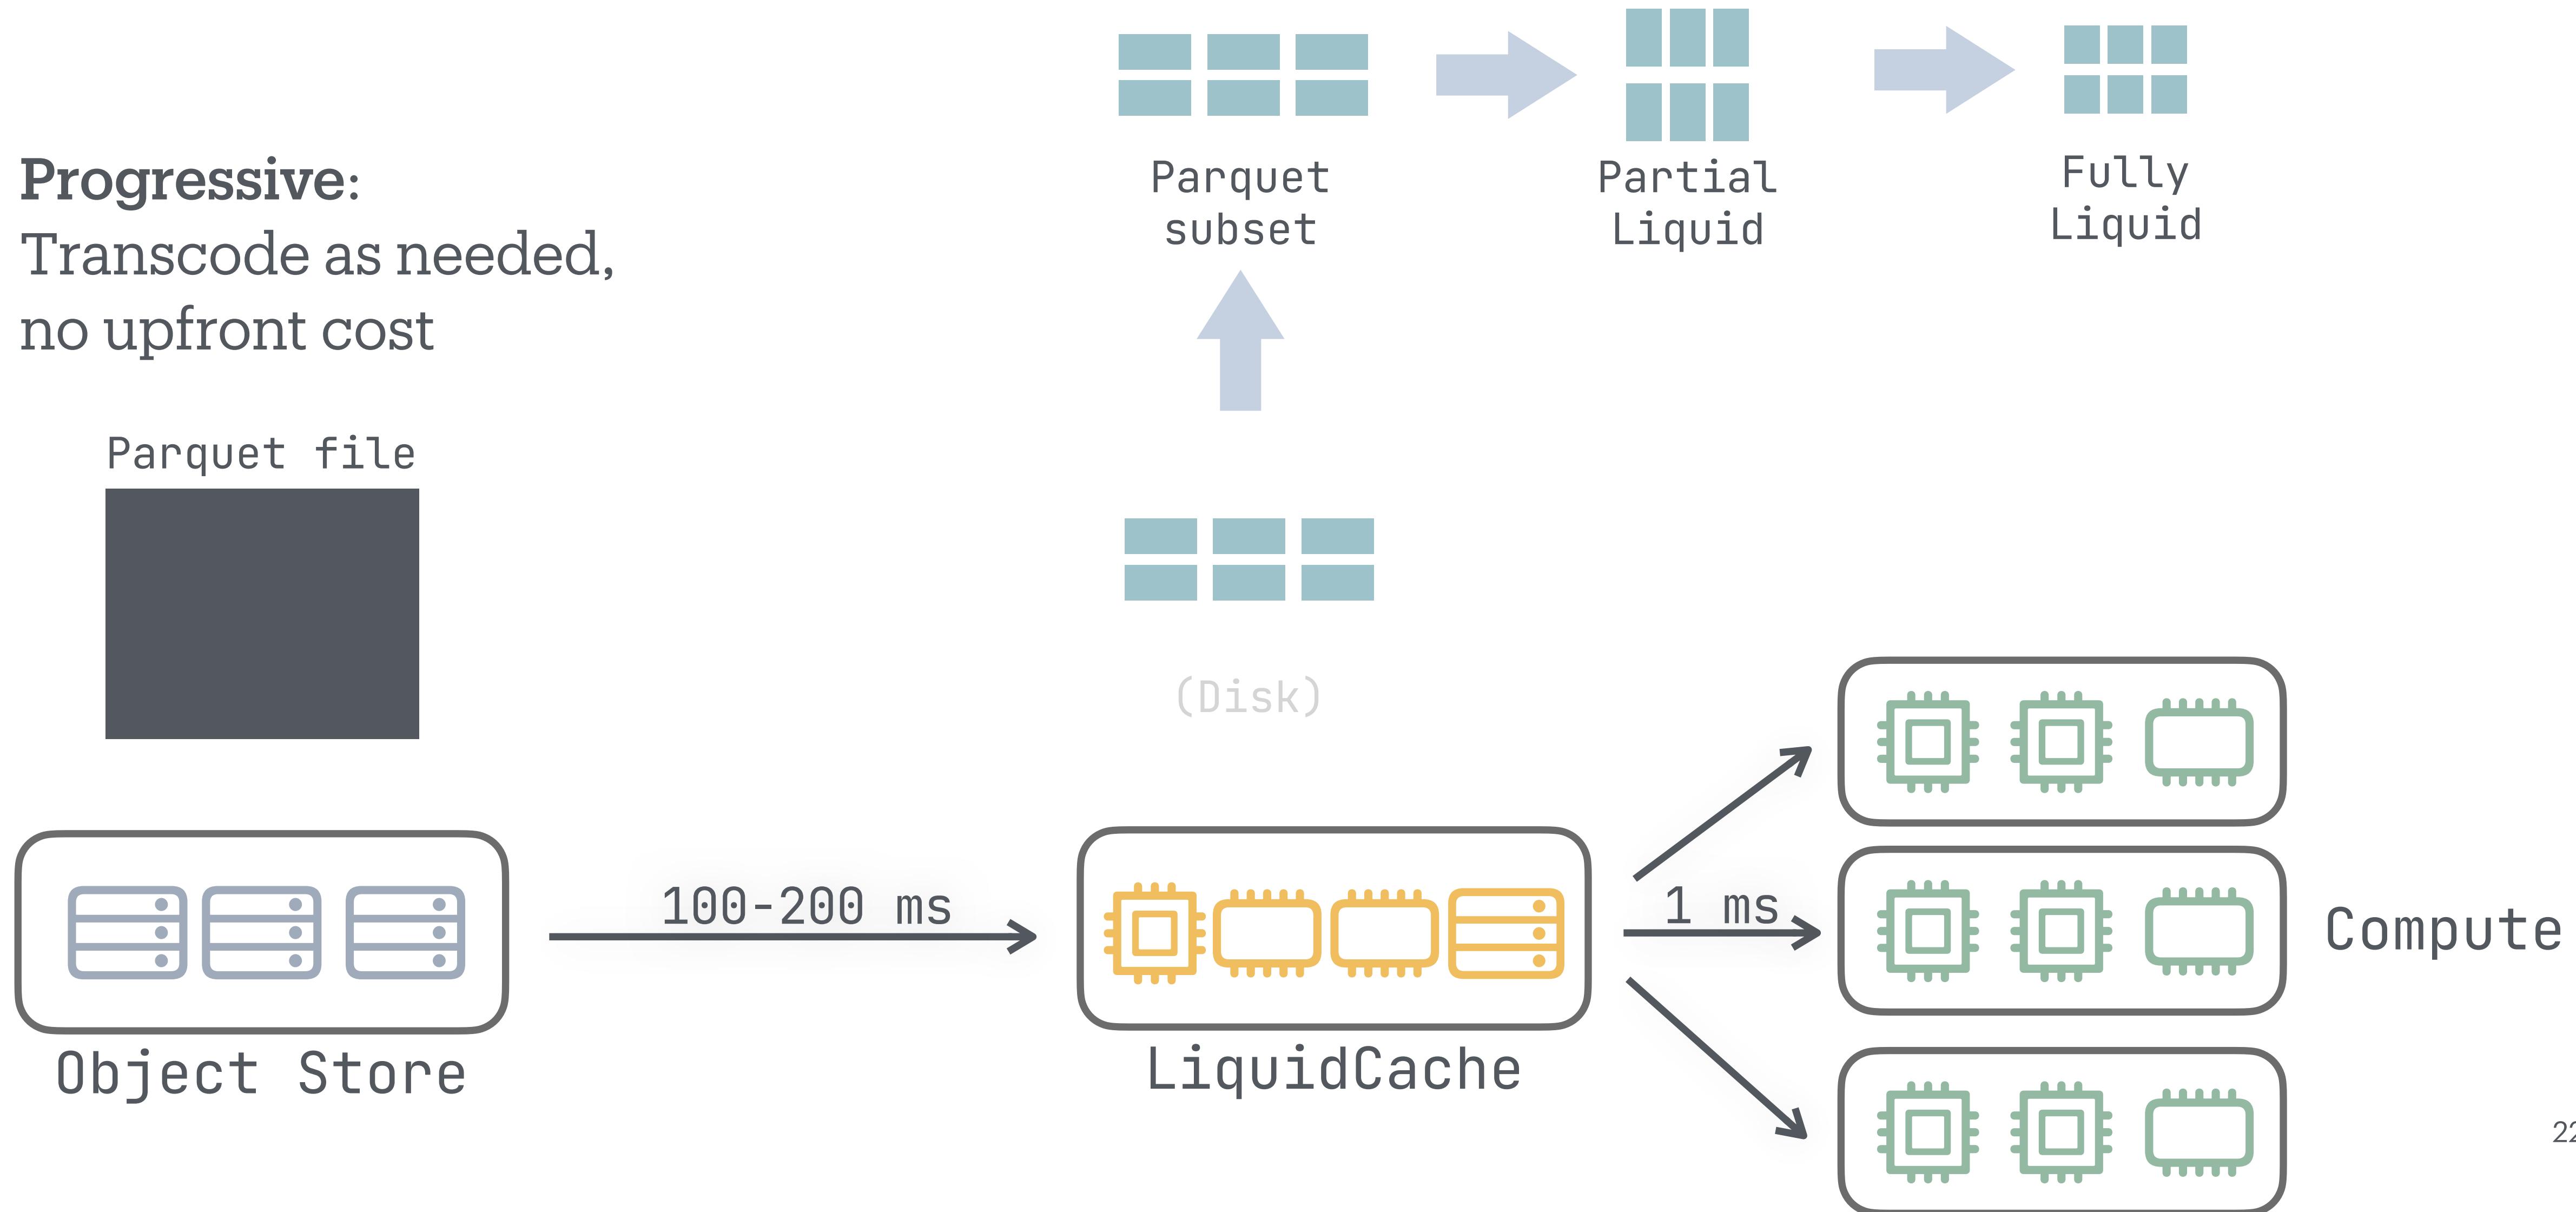
10x { cost | latency } reduction



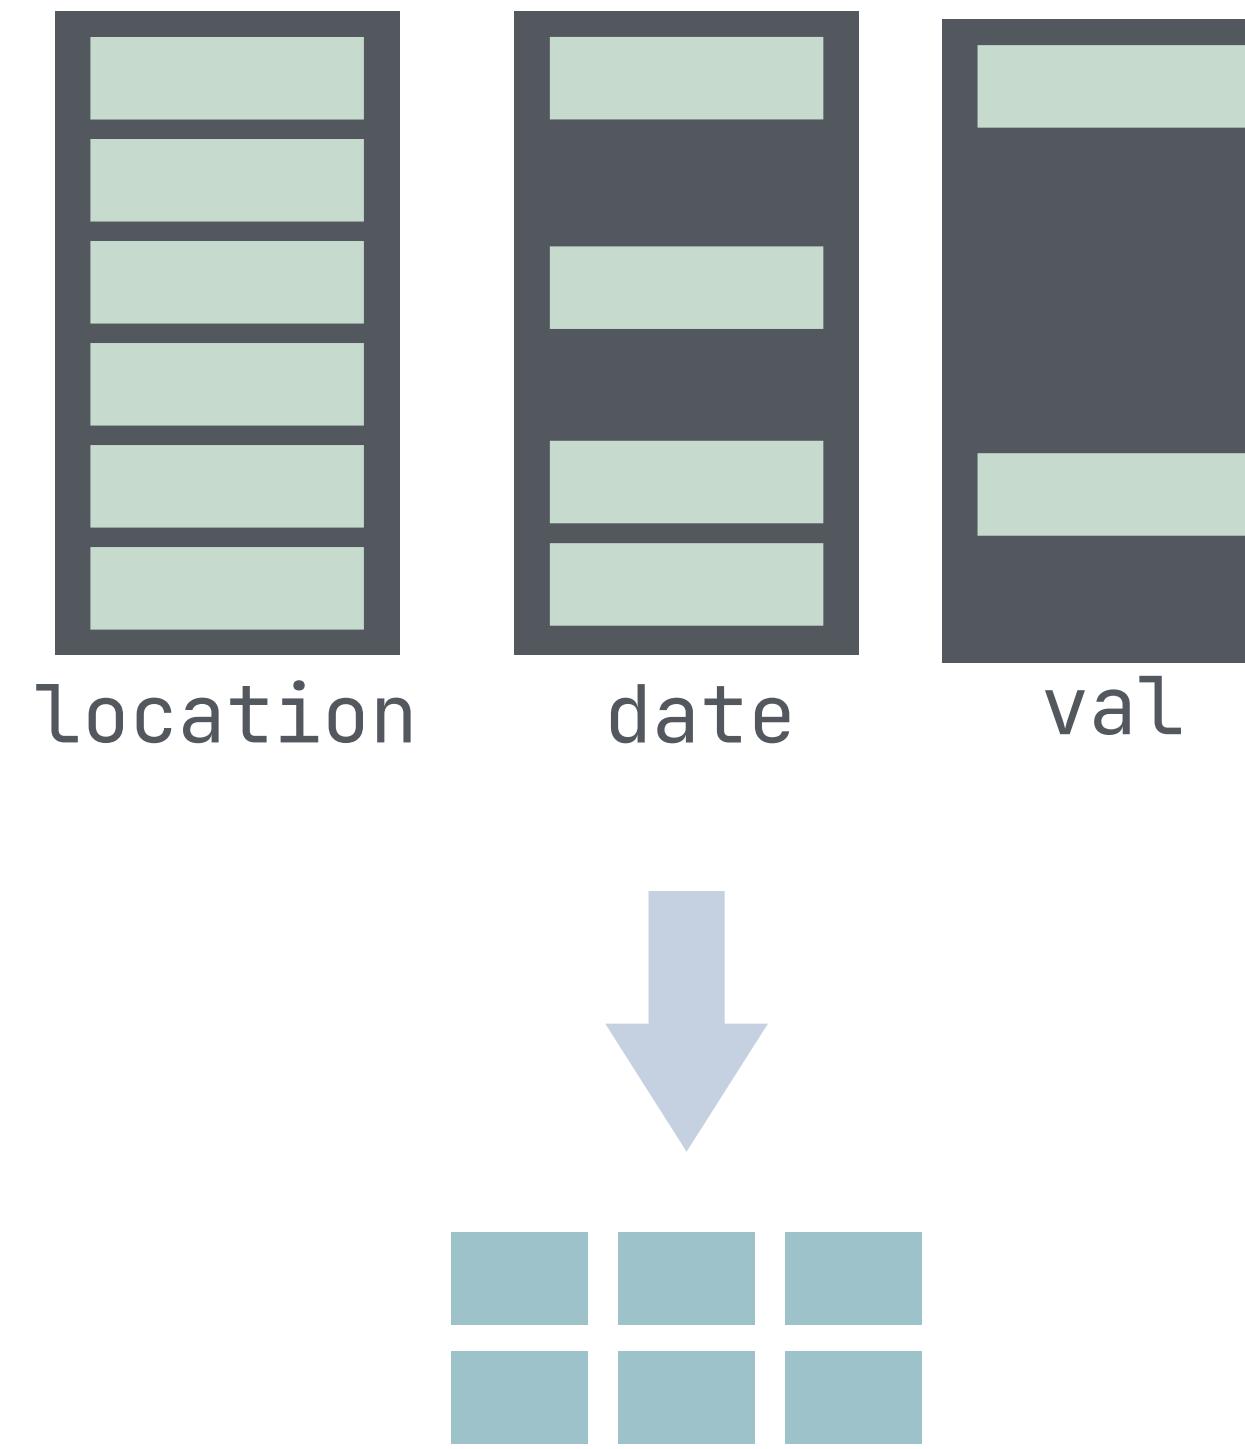
Decoding cost: close to theoretical optimal

Progressive transcoding

Progressive:
Transcode as needed,
no upfront cost

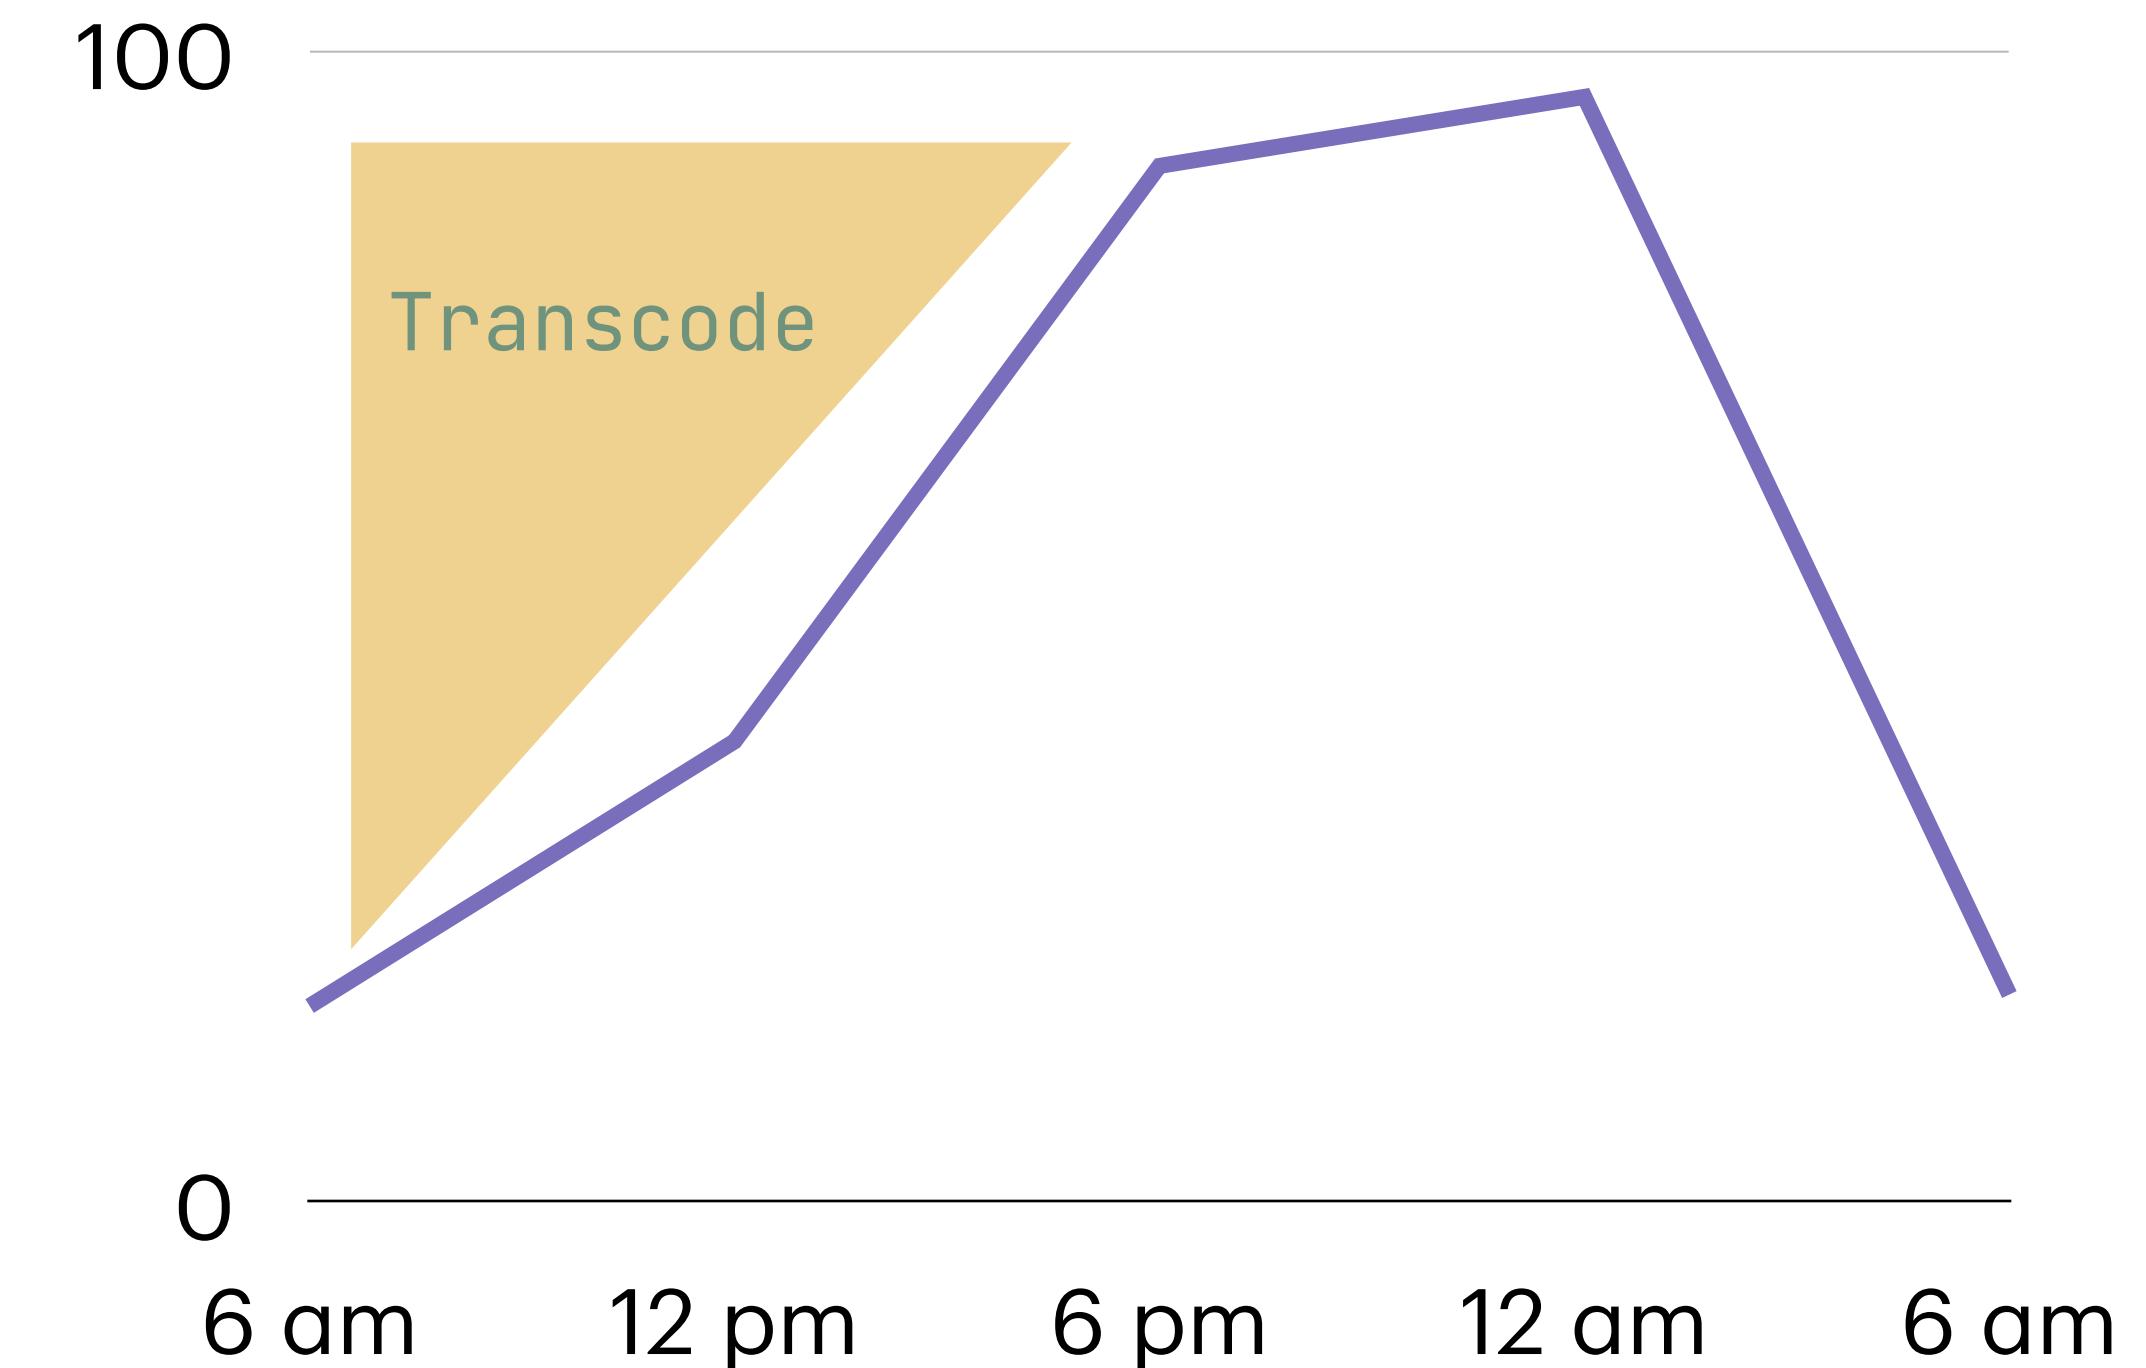


Selective transcoding

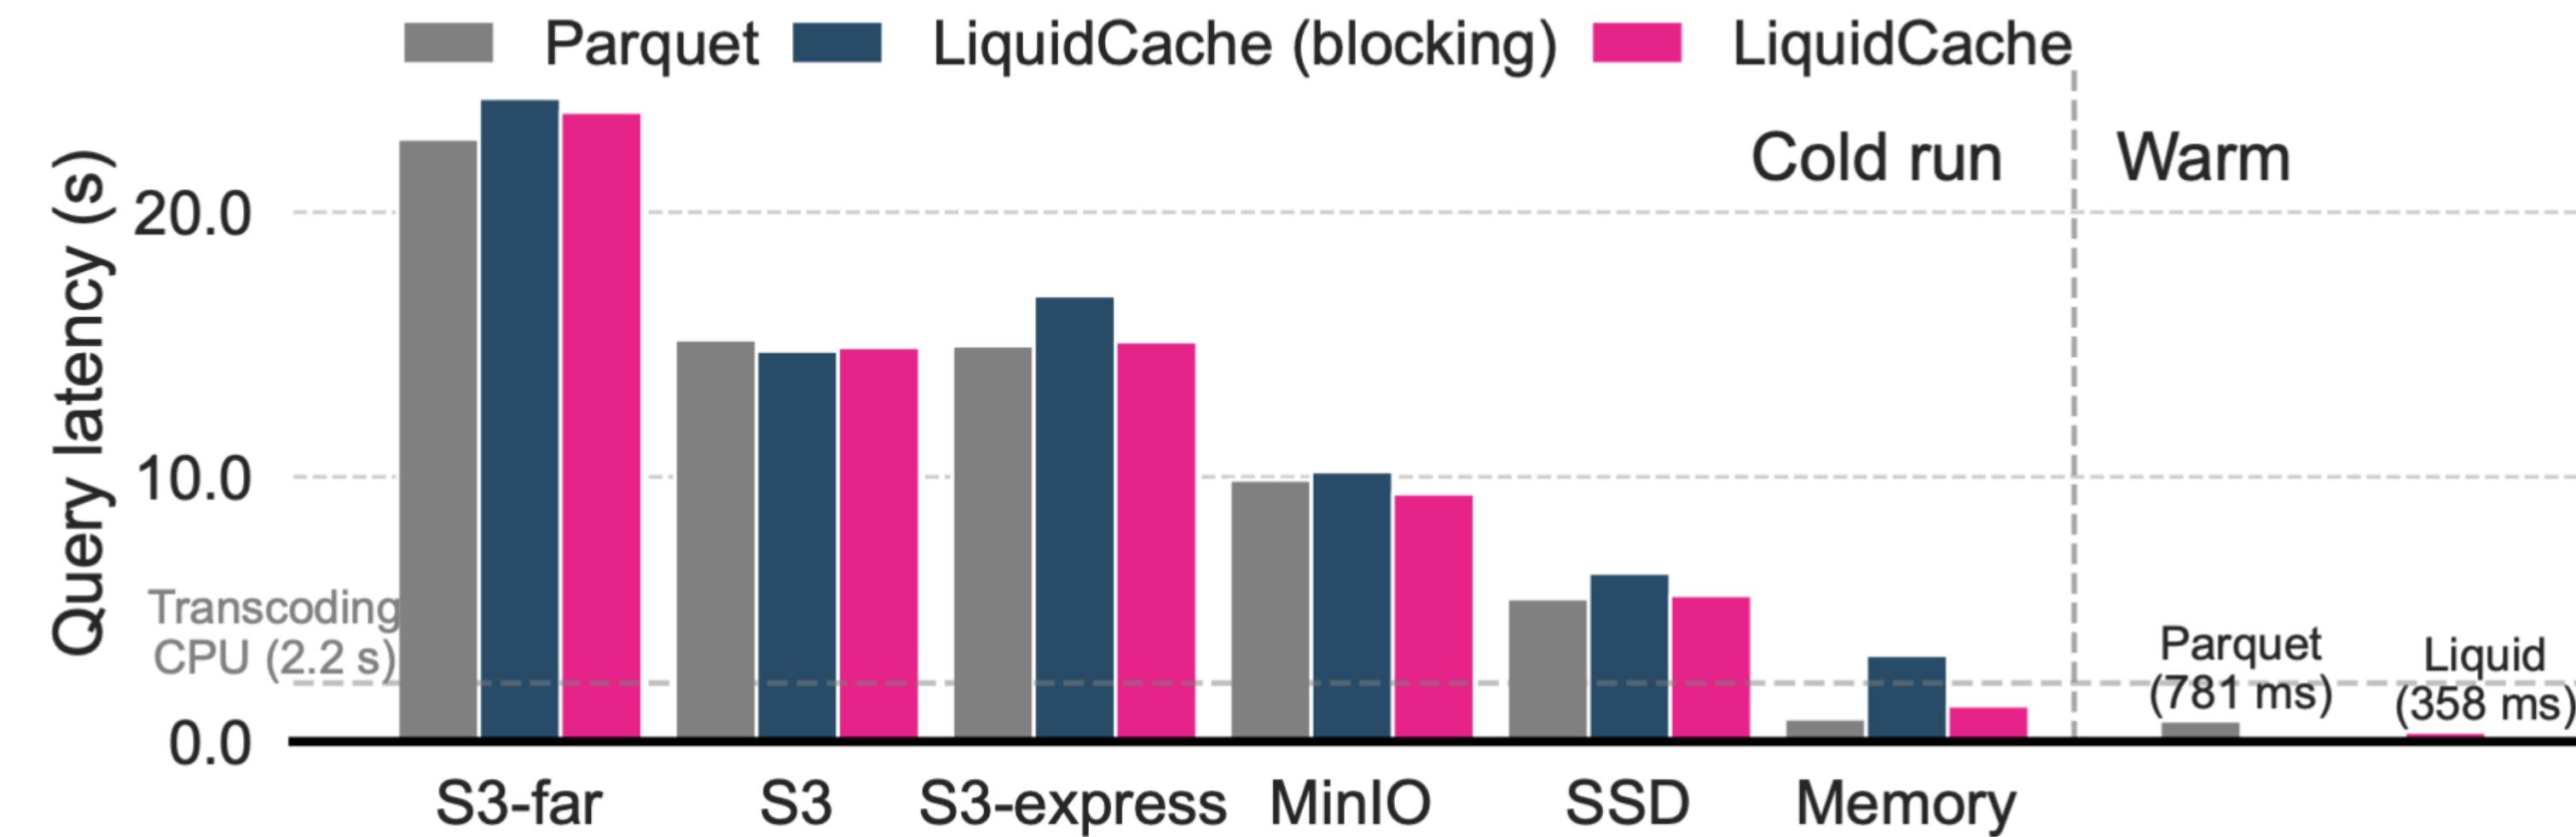


Selective:
Transcode only touched data

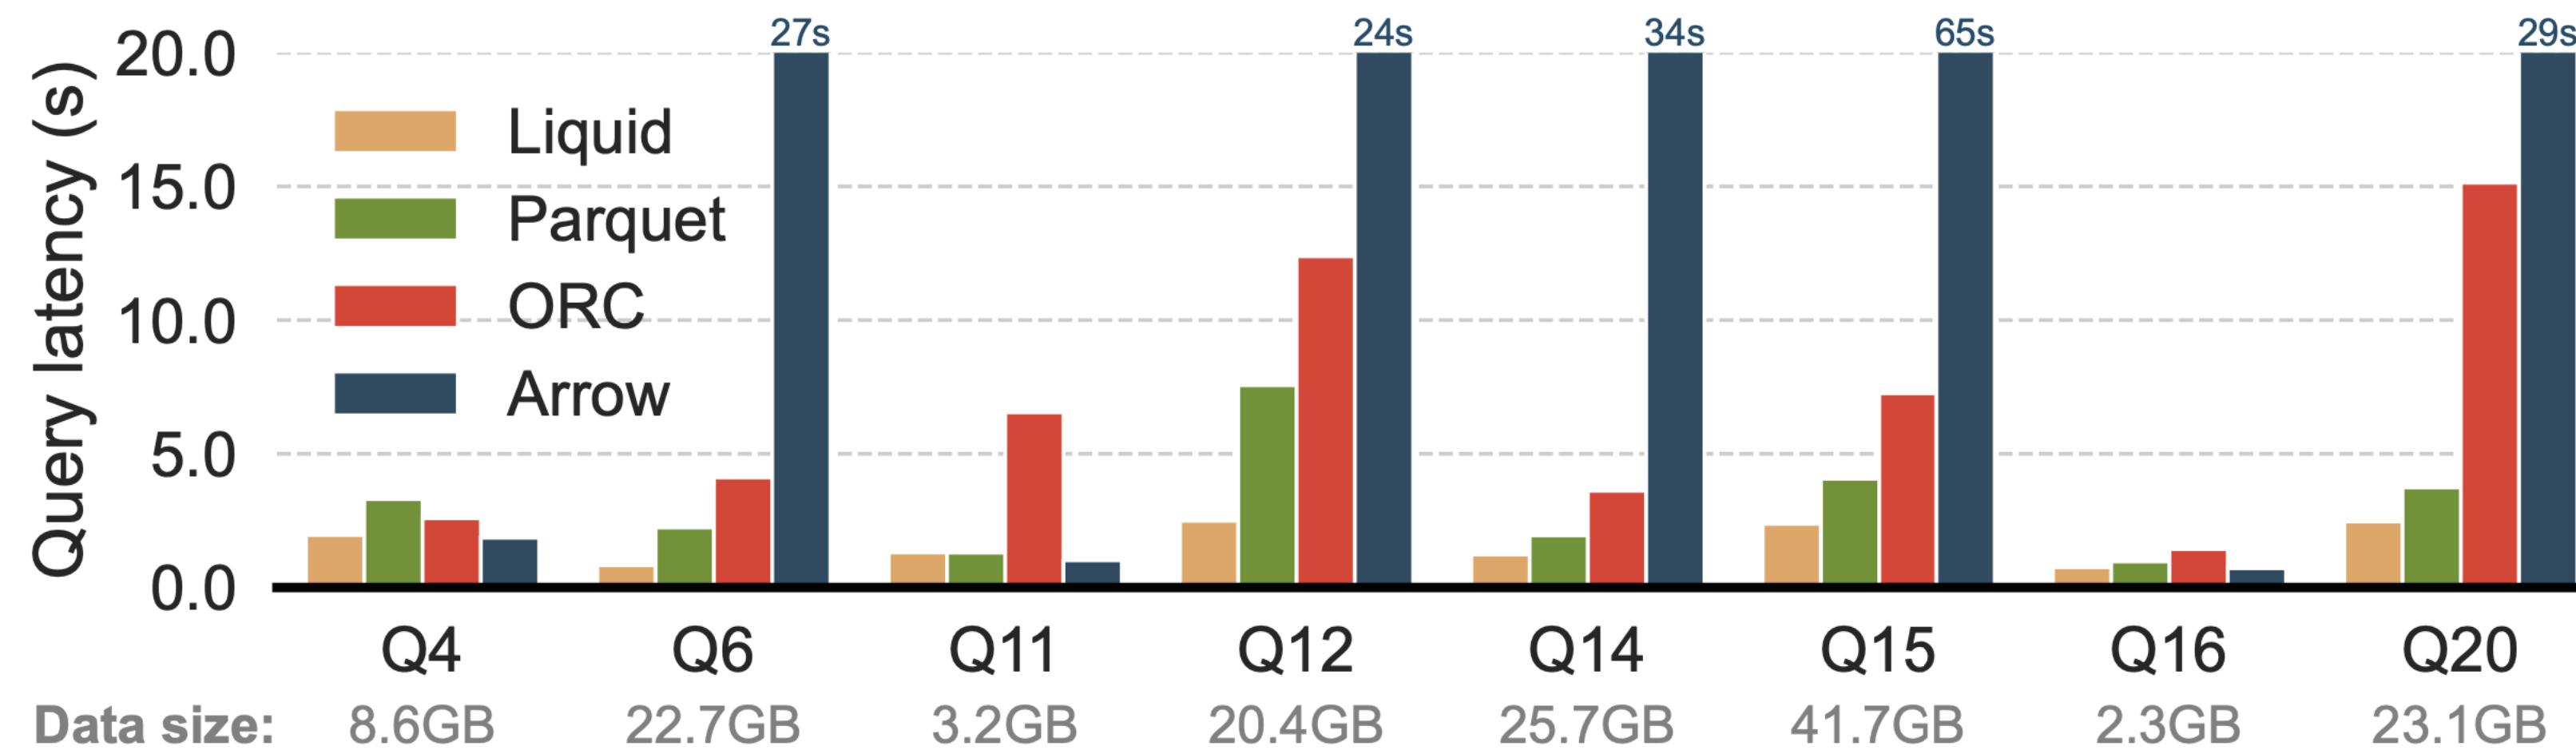
Background transcoding



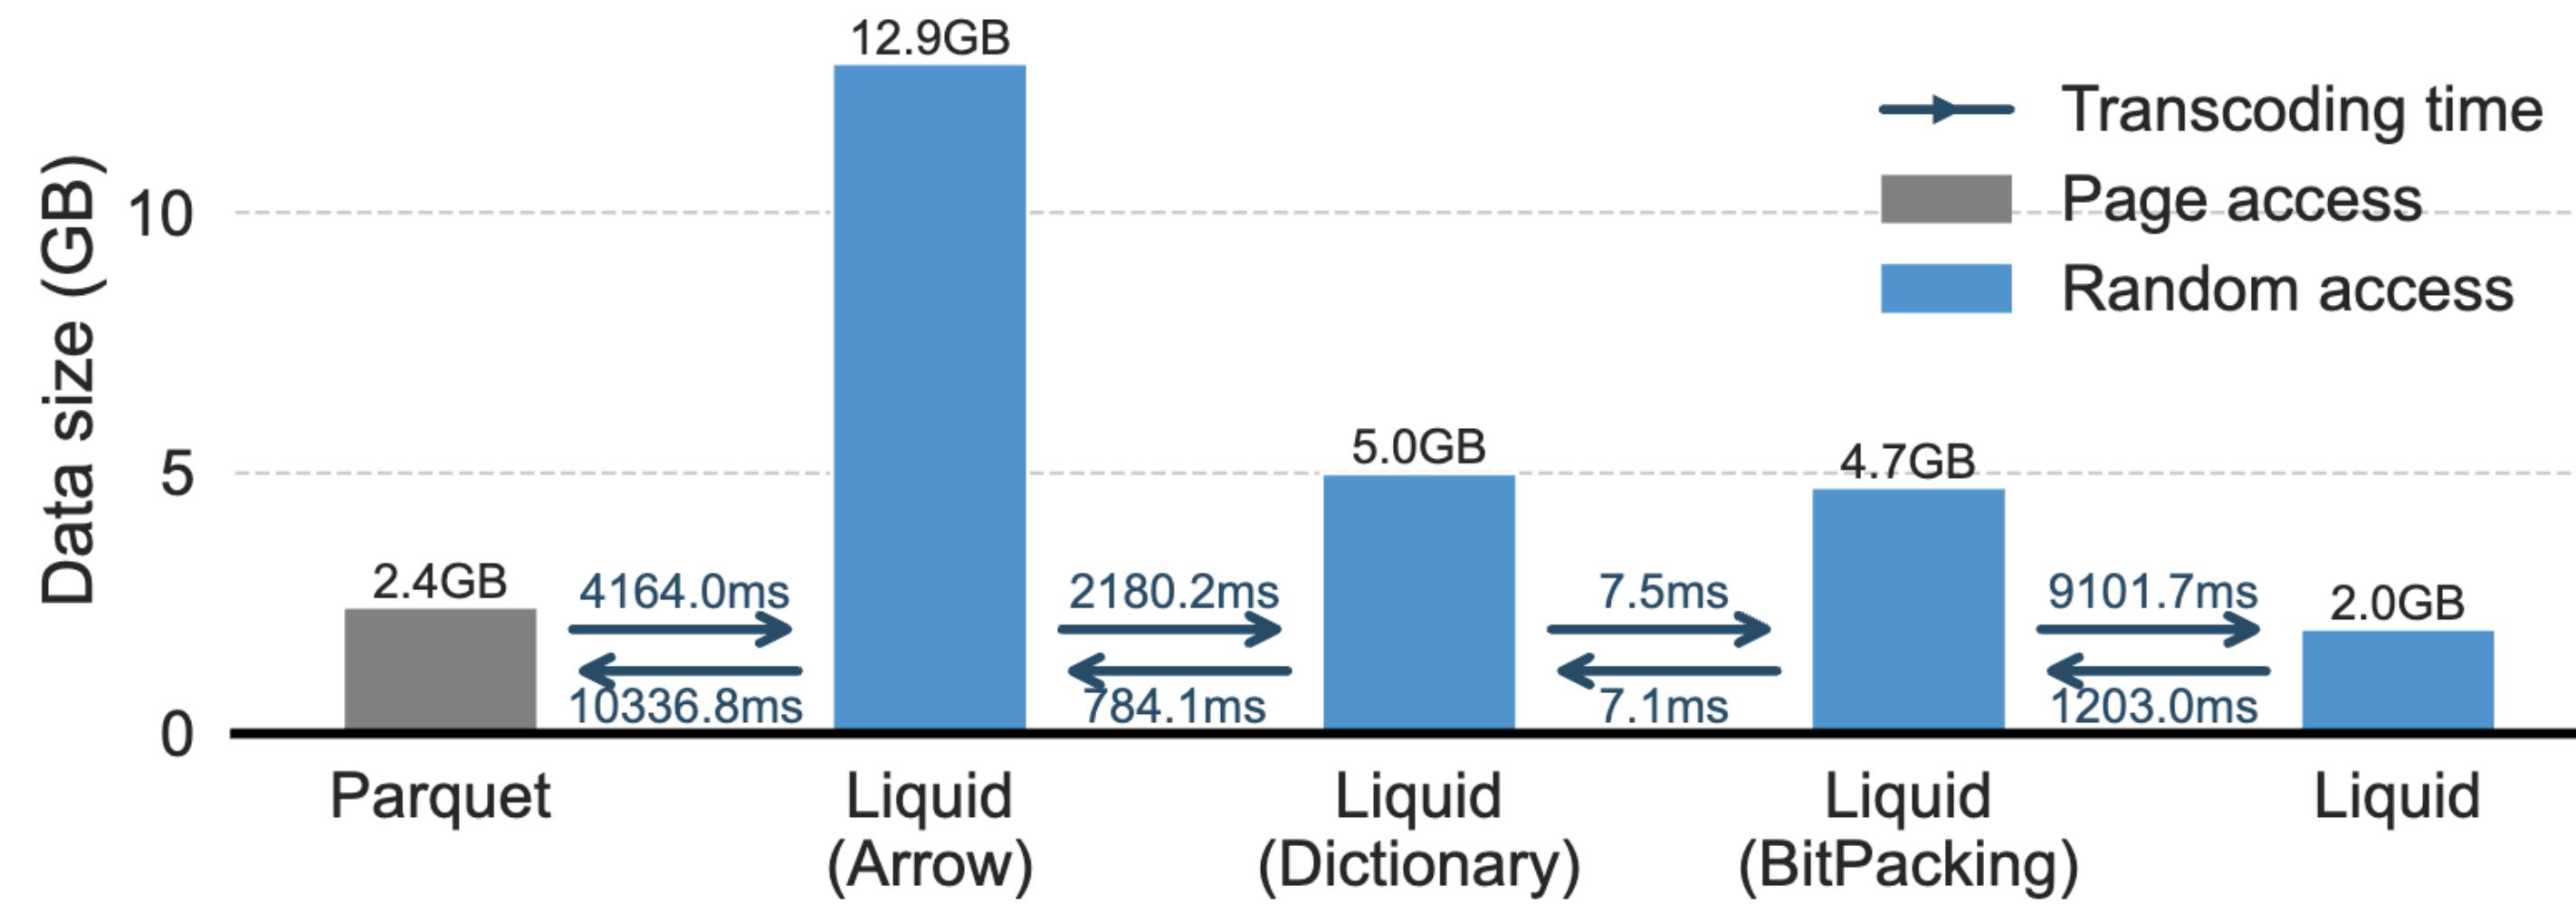
Asynchronous:
Transcode when less busy



Transcoding cost: hide by IO



With same memory: 10x lower latency



Compression ratio: comparable to Parquet

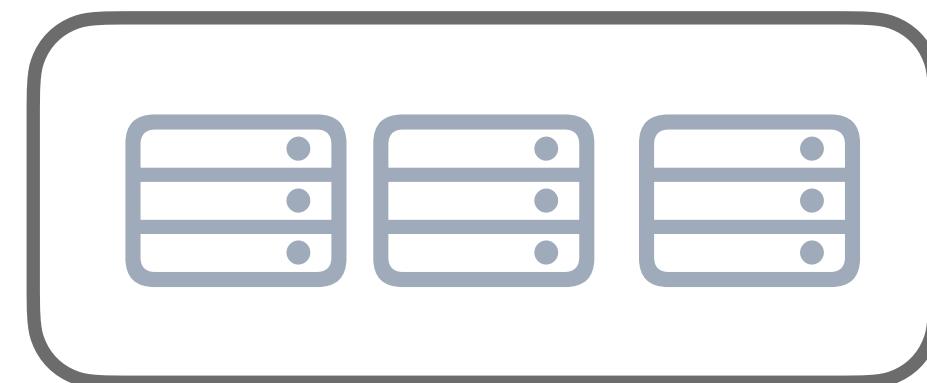


LiquidCache

10x { cost | latency } reduction

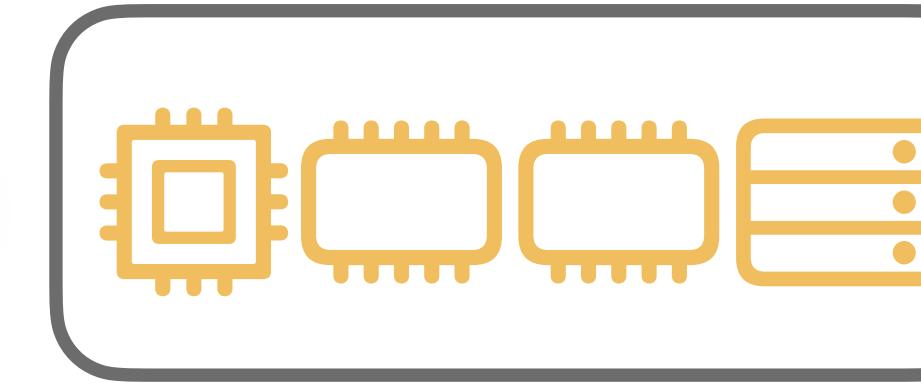


[https://github.com/
XiangpengHao/liquid-cache](https://github.com/XiangpengHao/liquid-cache)



Object Store

100-200 ms



LiquidCache

